

**A MAJOR PROJECT
REPORT
ON
Smart TB Surveillance: Integrating AI for Early Diagnosis and Spread
Monitoring
BACHELOR OF TECHNOLOGY
In
Computer Science & Engineering (AIML)
Submitted By
Rolika Agarwal (2101921530127)
Shivi Arora (2101921530153)
Srishti Singh (2101921530162)
Sweksha Dixit (2101921530170)**

**Under the Supervision of
Mr. Anand Kumar Yadav**

**G.L. BAJAJ INSTITUTE OF TECHNOLOGY &
MANAGEMENT, GREATER NOIDA**



**Affiliated to
DR. APJ ABDUL KALAM TECHNICAL UNIVERSITY,
LUCKNOW**



2024-2025

Declaration

We hereby declare that the project work presented in this report entitled "**Smart TB Surveillance: Integrating AI for Early Diagnosis and Spread Monitoring**", in partial fulfillment of the requirement for the award of the degree of Bachelor of Technology in Computer Science & Engineering (AIML) submitted to A.P.J. Abdul Kalam Technical University, Lucknow, is based on our own work carried out at Department of Computer Science & Engineering (AIML), G.L. Bajaj Institute of Technology & Management, Greater Noida. The work contained in the report is true and original to the best of our knowledge and project work reported in this report has not been submitted by us for award of any other degree or diploma in any other institute / college/university.

Signature:

Name:

Roll No:

Signature:

Name:

Roll No:

Signature:

Name:

Roll No:

Signature:

Name:

Roll No:

Date:

Place: Greater Noida

Certificate

This is to certify that the Project report entitled “ **Smart TB Surveillance: Integrating AI for Early Diagnosis and Spread Monitoring**” done by **Rolika Agarwal (2101921530127), Shivi Arora (2101921530153), Srishti Singh (2101921530162) and Sweksha Dixit (2101921530170)** is an original work carried out by them in Department of Computer Science & Engineering (AIML), G.L. Bajaj Institute of Technology & Management, Greater Noida under my supervision . The matter embodied in this project work has not been submitted earlier for the award of any degree or diploma to the best of my knowledge and belief.

Date:

Supervisor's Signature

Mr. Anand Kumar Yadav

HOD's Signature

Dr. Santosh Kumar Srivastava

Acknowledgement

We would like to express our sincere gratitude to everyone who contributed to the successful completion of this project, *Smart TB Surveillance: Integrating AI for Early Diagnosis and Spread Monitoring*.

First and foremost, we extend our heartfelt thanks to our project supervisor, Mr. Anand Kumar Yadav, for their invaluable guidance, support, and encouragement throughout this endeavour. Their expertise and insights were instrumental in shaping the direction of our work.

We pay special thanks to our Head of the Department Dr. Santosh Kumar Srivastava who has been always present as a support and help us in all possible way during this project.

We are grateful to GL Bajaj Institute of Technology and Management for providing the resources and infrastructure necessary for conducting this research. The access to computational facilities and datasets significantly contributed to the development of this project.

Lastly, we extend our deepest gratitude to our families, friends, and colleagues for their unwavering support and encouragement, which motivated us to complete this project successfully.

ABSTRACT

Tuberculosis (TB) remains a noteworthy worldwide wellbeing obstacle, particularly in low-resource settings where convenience and exact determination is regularly restricted. This project proposes an automated and efficient approach to Pulmonary TB detection by utilizing deep learning models trained on chest X-ray images. The workflow begins with preprocessing and converting DICOM images into PNG format, followed by an 80:10:10 split of the dataset for training, validation, and testing. Numerous pioneering convolutional neural network (CNN) architectures, including ResNet50V2, InceptionV3, and Xception, were explored and the final model was optimized using Keras Tuner for hyperparameter tuning. Early stopping method was employed to enhance generalizability and reduce overfitting.

Model performance was assessed utilizing key measurements such as AUC-ROC , F1-Score, Accuracy and PR-AUC. The ResNet50V2 model, tuned with optimal hyperparameters, demonstrated superior classification results. The final model was deployed using the Streamlit framework, enabling users to make use of the TB detection system using a simple web interface. Although Grad-CAM visualization was not incorporated, the deployment provides a functional and accessible diagnostic aid for early TB screening. The outcomes indicate promising potential for integrating AI-powered TB detection into clinical workflows. Future enhancements could include ensemble learning, Grad-CAM-based interpretability, and the expansion of the model to detect other chest-related diseases.

TABLE OF CONTENT

Declaration	(ii)
Certificate	(iii)
Acknowledgement	(iv)
Abstract	(v)
Table of Content.....	(vi)
List of Figures	(vii)
List of Tables.....	(viii)
Chapter 1. Introduction	1-7
1.1 Background & Motivation.....	1
1.2 Problem Statement & Objectives	2
1.3 Benefits of research.....	6
1.4 Limitation of research.....	6
Chapter 2. Literature Survey	8-21
2.1 Introduction	8
2.2 Literature Review.....	8
2.3 Inferences Drawn from Literature Review.....	20
Chapter 3. Proposed Work.....	22-27
3.1 Introduction	22
3.2 Proposed Work	22
Chapter 4. Methodology.....	28-45
4.1 Introduction	28
4.2 Implementation Strategy	28
4.3 Methodology Steps.....	30
4.4 Tools & Technologies.....	44
Chapter 5. Result & Discussion	46-59
5.1 Result	46
5.2 Discussion.....	55
Chapter 7. Conclusion and Future Scope.....	60-61
7.1 Conclusion.....	60
7.2 Future Scope	61

References

Plagiarism Report

LIST OF FIGURES

S.No	Figure No.	Description	Page No.
1	Figure 3.1	Workflow Diagram of the Proposed TB Detection System	24
2	Figure 4.1	Flowchart of the TB Detection Methodology Pipeline	29
3	Figure 4.2	Code Snippet of DICOM Conversion	30
4	Figure 4.3	Snippet of Source Folders	30
5	Figure 4.4	Snippet of DICOM Conversion Terminal	31
6	Figure 4.5	Code Snippet of Data Splitting	32
7	Figure 4.6	Snippet of Data Splitting terminal	33
8	Figure 4.7	Code Snippet of K-Fold Cross-Validation Function	34
9	Figure 4.8	Code Snippet of Dataset tensor creation	34
10	Figure 4.9	Code Snippet for Fold 3 of EfficientNetB0 training	35
11	Figure 4.10	Code Snippet for Fold 2 of MobileNetV2 training	36
12	Figure. 4.11	Code Snippet for Fold 4 of Resnet50 training	36
13	Figure. 4.12	Code Snippet for Fold 3 of DenseNet121 training	37
14	Figure. 4.13	Code Snippet for Fold 3 of InceptionV3 training	37
15	Figure. 4.14	Code Snippet for Fold 3 of Xception training	38
16	Figure. 4.15	Snippet of Model Comparison terminal	38
17	Figure. 4.16	Code Snippet of Hyperband definition	39
18	Figure. 4.17	Code Snippet of Loading Datasets	39
19	Figure. 4.18	Snippet of Hypertuning terminal	40
20	Figure. 4.19	Code Snippet of building final model	41
21	Figure. 4.20	Snippet of best hyperparameters	41
22	Figure. 4.21	Snippet of final model training terminal	42
23	Figure. 4.22	Snippet of final model test results	42
24	Figure. 4.23	Code Snippet of Model Deployment function	43
25	Figure. 4.24	Snippet of Web UI	44
26	Figure. 5.1	Model Performance Comparison Heatmap	48
27	Figure. 5.2	PR-AUC Comparison	49
28	Figure. 5.3	ROC-AUC Comparison	50
29	Figure. 5.4	ROC-AUC and PR AUC Comparison	50
30	Figure. 5.5	ROC Curve showing high true positive rate with ResNet50V2	51
31	Figure. 5.6	Mean ROC Curve	51
32	Figure. 5.7	Validation accuracy/loss curve	52
33	Figure. 5.8	Train Vs Test Metrics across epochs	53
34	Figure. 5.9	Confusion matrix of ResNet50V2 classification distribution	53
35	Figure. 5.10	Precision-Recall Curve for the Final Model on the Test Set	54

LIST OF TABLES

S.no	Table No.	Description	Page No.
1	Table 5.1	Model Summary Report	47

Chapter 1

1. Introduction

1.1 Background Information:

Tuberculosis (TB) is a contagious disease wrought by the bacterium known as *Mycobacterium tuberculosis*. It predominantly affects the lungs (pulmonary TB) but can also impact other parts of the body, including the spine, kidney and brain (extra-pulmonary TB). TB transmits through the air when people with active pulmonary TB sneeze, spit or cough, expelling germs into the air. Individuals then become infected by inhaling these germs.

In spite of being avertible and treatable, TB continues to be the one of the top contagious diseases resulting in deaths across the globe. According to the World Health Organization (WHO), approximately 10.6 million people fell sick with TB globally in 2022, and 1.3 million succumbed to the disease. Another thing to make note of is that the impact of disease is felt disproportionately high by the low- and middle-income countries. Symptoms of pulmonary TB can include a fever, chest pains, night sweats, relentless cough (sometimes with blood) and weight loss. Early and accurate diagnosis is paramount not only for timely patient treatment and improved health outcomes but also for preventing further transmission within communities. However, diagnosing TB, especially in resource-limited settings, can be challenging due to non-specific symptoms and the limitations of current diagnostic methods.

Existing Solutions for TB Detection:

Determination of Pulmonary Tuberculosis generally includes a combination of clinical assessment, microbiological tests, and imaging techniques. Each method has its strengths and limitations, which collectively underscore the need for more efficient and universally accessible diagnostic tools.

Clinical Checkup: This encompasses determining a patient's symptoms, physical signs and medical history. While crucial for initial suspicion, clinical examination alone is often insufficient for definitive diagnosis due to the non-specific nature of TB symptoms.

Sputum Smear Microscopy: This is a commonly used, fast, and inexpensive method for detecting acid-fast bacilli (AFB) in sputum samples. It identifies highly contagious individuals. However, its sensitivity is relatively low, particularly in patients with smear-negative TB (where bacterial load is low), extra-pulmonary TB, or in children. It also requires qualified and experienced microscopists and laboratory infrastructure.

Mycobacterial Culture: This method involves growing *Mycobacterium tuberculosis* in a laboratory. It is regarded as the gold standard for TB diagnosis due to its high sensitivity, allowing for the detection of low bacterial concentrations and drug susceptibility testing. However, cultures are slow, taking several weeks for results, which can delay treatment and increase transmission.

Molecular Tests (Nucleic Acid Amplification Tests - NAATs): Technologies like Xpert MTB/RIF and TrueNat are rapid and highly sensitive molecular assays that detect *Mycobacterium tuberculosis* DNA and genetic mutations linked to drug resistance. While highly effective, these tests are significantly more expensive than microscopy and require specialized equipment and trained personnel, limiting their widespread accessibility in all settings.

Chest Radiography (CXR): Chest X-rays are generally utilized as a screening instrument and for detecting pulmonary TB. They can expose characteristic irregularities like pleural effusions, cavities, nodules, and infiltrates. CXR is non-invasive and widely available. However, its interpretation is subjective and highly dependent on the experience of the radiologist, leading to potential inter-observer variability. Furthermore, certain radiological findings can be ambiguous or mimic other lung diseases, making definitive diagnosis challenging, especially in high-volume screening programs.

The limitations inherent in these existing diagnostic methods—including varying levels of sensitivity and specificity, the time required for results, high costs, and dependence on specialized infrastructure or human expertise—collectively highlight a significant diagnostic gap. This gap motivates the exploration of alternative, complementary technologies that can offer faster, more objective, and accessible diagnostic assistance, particularly in resource-constrained environments.

Tuberculosis (TB) remains to be a worldwide health malady challenge, especially in less-and average-income countries where access to healthcare infrastructure and expert radiologists is limited. As reported by the World Health Organization (WHO), TB is one of the leading infectious disease killers, causing over 1.3 million deaths annually. Early diagnosis plays a critical role in effective treatment and containment, but traditional diagnostic techniques such as sputum tests and manual analysis of chest X-rays are time-intensive, prone to human error, and often unavailable in remote areas.

Chest X-ray (CXR) imaging is a widely used, non-invasive, and cost-effective tool for TB screening. However, interpreting CXRs for TB diagnosis requires expert radiologists, and there exists substantial inter-reader variability. Moreover, the growing patient burden in many countries creates a bottleneck for timely screening and diagnosis.

To address this challenge, our work proposes a smart surveillance system that utilizes deep learning techniques—particularly Convolutional Neural Networks (CNNs)—for automated detection of tuberculosis from chest X-ray images. Deep learning has shown remarkable success in image classification tasks, and pre-trained CNN architectures such as ResNet50, EfficientNetB0, and MobileNetV2 offer powerful tools for medical image analysis through transfer learning.

In our project, we evaluated six state-of-the-art CNN models on a well-curated TB chest X-ray dataset. The models were trained using TensorFlow and Keras, with extensive hyperparameter tuning applied to ResNet50 using Keras Tuner. Evaluation metrics such as accuracy, F1-score, ROC AUC, and PR AUC were used to assess each model's performance. Among them, ResNet50 emerged as the best performer, achieving an average accuracy of **89.95%**, F1-score of **89.79%**, and ROC AUC of **96.43%**, demonstrating both

reliability and generalization capability.

This project is inspired by the need for a flexible, consistent, and resource-efficient solution to support clinical decision-making and large-scale TB screening efforts. By deploying such AI-based systems in diagnostic workflows, healthcare providers can reduce diagnostic delays, improve patient outcomes, and extend services to under-resourced regions.

1.2 Problem Statement and Objectives:

Tuberculosis (TB) remains to be a principal driver of morbidity and mortality worldwide, especially in developing countries where access to timely and accurate diagnostic tools is limited. Timely identification of TB is imperative for effective treatment and obviating disease transmission. However, traditional methods for diagnosing TB, such as sputum smear microscopy and culture tests, are time-consuming, resource-intensive, and often inaccessible in low-resource settings.

Additionally, the reliance on manual interpretation of diagnostic images, such as chest X-rays, introduces the risk of human error and variability in diagnosis. This poses a significant challenge in areas with limited availability of skilled radiologists and healthcare professionals.

The core problem lies in the absence of scalable, automated, and reliable TB detection mechanisms that can match or exceed human-level performance, especially in rural or under-resourced healthcare systems. While machine learning offers a promising path forward, building an accurate model for TB detection involves several technical challenges:

- Selecting an optimal CNN architecture that balances accuracy, inference speed, and training time.
- Fine-tuning hyperparameters (e.g., learning rate, dropout rate, dense units) to prevent overfitting and underfitting.
- Handling dataset imbalances, variance in image quality, and potential overfitting during training.
- Ensuring reproducibility and generalizability across unseen data.

To address these issues, this study formulates the following research problem:

“How can deep learning models, particularly optimized Convolutional Neural

Networks, be effectively designed, trained, and evaluated for accurate and efficient tuberculosis detection from chest X-ray images?”

This leads to several sub-problems tackled during the project:

- **Model Selection:** A comparative study was performed using six popular CNN architectures—EfficientNetB0, MobileNetV2, ResNet50, DenseNet121, InceptionV3, and Xception—each pre-trained on ImageNet.
- **Hyperparameter Optimization:** Keras Tuner's Hyperband strategy was employed to explore the best configuration of ResNet50 parameters, aiming to maximize the validation AUC.
- **Performance Evaluation** The best model (ResNet50) was evaluated on key metrics—Accuracy, F1-score, ROC AUC, PR AUC—and analyzed for training efficiency.
- **Generalization and Reliability:** The best model (ResNet50) was validated with plots showing tuning trends and metric stability, ensuring robust performance across multiple trials.

The comparative results showed that **ResNet50**, with tuned hyperparameters, achieved the best trade-off between predictive performance and computational cost, outperforming others on key metrics (e.g., AUC: **0.9643**, F1: **0.8979**, PR AUC: **0.9673**).

This project, therefore, sets out to solve the real-world problem of automating TB detection using AI, aiming to deliver a deployable, smart TB surveillance system that is both accurate and scalable.

The main objectives of this project are:

1. To implement a process for converting chest X-ray images from the DICOM format to a standard image format (PNG) suitable for deep learning processing.
2. To evaluate the performance of multiple pre-trained Convolutional Neural Network (CNN) architectures using K-Fold Cross-Validation, to obtain reliable estimates of model performance and use transfer learning as feature extractors for TB detection.
3. To perform hyperparameter tuning for a selected model architecture to optimize its performance and train a final deep learning model on the combined training and validation

data using the optimized hyper parameters.

4. To develop a simple application for deploying the trained model to allow for user interaction and practical inference.

1.3 Benefits of research

- **Scalability:** Enables mass TB screening programs using digital infrastructure.
- **Improved diagnostic efficiency:** Through automation, the system can rapidly screen images, potentially reducing the time required for initial diagnosis.
- **Enhanced diagnostic accessibility:** The system can extend diagnostic capabilities to remote or underserved areas lacking specialist radiologists.
- **Increased objectivity and reduced subjectivity:** Automated analysis minimizes human interpretation variability, leading to more consistent diagnoses.
- **Creation of a foundational pipeline:** The developed workflow for DICOM processing and deep learning application serves as a robust baseline for future medical image analysis projects.
- **Making DICOM data accessible:** Providing a pipeline to convert DICOM data into standard image formats facilitates broader research and development using conventional image processing and deep learning tools.
- **Healthcare Impact:** Supports public health programs, especially in resource-poor regions.

1.4 Limitation of Research

While the study successfully demonstrates the effectiveness of deep learning models, particularly ResNet50, in detecting tuberculosis from chest X-ray images, several limitations constrain the scope and generalizability of the findings:

- **Dataset Constraints:**

The models were trained and evaluated on a curated dataset composed of publicly available TB chest X-ray images. Although efforts were made to improve diversity by combining multiple sources (Kaggle, Mendeley, IN-CXR), the dataset may not represent the full spectrum of radiographic variations observed across different demographics, geographic

regions, and equipment types. As a result, the models may exhibit degraded performance when applied to unseen clinical environments.

- **Computation and Training Variability:**

While hyperparameter tuning using Keras Tuner (Hyperband) improved model performance, it introduced significant computational overhead. Some models such as DenseNet121 exhibited high variance in training time (± 2553.78 seconds), indicating sensitivity to hardware load, batch sizes, or optimizer dynamics, which may affect reproducibility on different systems.

- **Limited Ensemble or Multimodal Integration:**

Although multiple CNN architectures were compared, the study did not explore ensemble techniques or integration with non-imaging features (e.g., patient history, lab tests), which could further enhance diagnostic accuracy and reliability in practical applications.

Chapter 2

2. Literature Review

The accelerated evolution of deep learning, particularly in computer vision, has opened new avenues for solving complex problems across various domains, including medical diagnostics. To ensure the successful development of a robust and effective deep learning system for Tuberculosis detection, it is imperative to establish a strong foundational understanding of the existing body of knowledge. This chapter provides a comprehensive review of relevant literature, exploring the foundational concepts and state-of-the-art advancements pertinent to this project. The primary aim is to establish a theoretical and practical groundwork by analyzing established methodologies, identifying key challenges, and understanding performance benchmarks achieved in similar research endeavors. Specifically, this literature survey delves into the application of Convolutional Neural Networks (CNNs) and transfer learning in medical image analysis, with a particular focus on automated Tuberculosis detection from chest X-rays. It also examines various prominent CNN architectures, explores robust evaluation methodologies suitable for medical imaging tasks, and reviews techniques for handling specialized medical image data formats like DICOM. The insights garnered from this review have directly informed the design and selection of the methods implemented in this project, as detailed in Chapter 3 (Proposed Work) and Chapter 4 (Methodology).

Paper 1: "Advances in Deep Learning for Tuberculosis Screening Using Chest X-rays: The Last 5 Years Review" by KC Santosh, et al.[1]

Abstract:

This paper reviews advancements in deep learning approaches for TB detection using chest X-rays over the last five years. The focus is on leveraging CNN architectures, transfer learning, and novel preprocessing methods to improve diagnostic accuracy.

Introduction:

The authors address the persistent global TB burden and the challenges associated with traditional diagnostic approaches. The review highlights how deep learning innovations are addressing gaps in TB diagnostics, particularly in underserved regions.

Background:

TB diagnostics through chest X-rays often suffer from variability due to radiologist expertise and resource availability. Recent advancements in AI and deep learning provide scalable solutions that are both efficient and reliable.

Methodology:

- The paper systematically reviews deep learning approaches applied to TB detection, focusing on techniques such as transfer learning, data augmentation, and ensemble methods.
- Case studies were analyzed to understand model generalizability across diverse datasets.
- Performance metrics like accuracy, sensitivity, and specificity were compared across studies.

Observation:

- Transfer learning significantly enhances model performance, particularly for limited datasets.
- Data preprocessing techniques, including augmentation and normalization, play a critical role in improving model robustness.
- Models achieved state-of-the-art performance when combined with external validation datasets.
-

Results:

The reviewed studies consistently demonstrated high accuracy (>90%) and low false negative rates, reinforcing the potential of deep learning in TB screening. However, challenges related to dataset quality, standardization, and real-world applicability persist.

Paper 2: "Chest X-Ray Analysis of Tuberculosis by Deep Learning with Segmentation and Augmentation" by S. Stirenko, et al.[2]

Abstract:

This paper introduces a deep learning-based method for TB detection from chest X-rays, incorporating image segmentation and data augmentation to enhance model accuracy and robustness.

Introduction:

The authors highlight the limitations of traditional diagnostic methods and emphasize the need for automated tools to support radiologists. The study explores the role of segmentation and augmentation in improving TB detection accuracy.

Background:

Segmentation techniques can isolate regions of interest in CXRs, reducing noise and improving model focus on diagnostically relevant areas. Augmentation addresses data scarcity and improves model generalizability.

Methodology:

- The authors implemented a CNN model with a preprocessing pipeline that included segmentation to isolate lung regions and augmentation to increase dataset diversity.
- Segmentation was performed using a U-Net architecture to enhance feature extraction.
- Augmentation techniques, such as rotation, scaling, and flipping, were applied to mitigate overfitting.

Observation:

- The segmentation step significantly reduced false positives by focusing the model on relevant areas of the image.
- Data augmentation improved model generalizability, particularly on unseen datasets.

Results:

The model achieved high performance, with an accuracy of 87% and improved sensitivity due to segmentation. The use of augmentation further enhanced the model's ability to handle variability in CXRs.

Paper 3: “AI for Detection of Tuberculosis: Implications for Global Health” by Hwang EJ, et al.[3]

Abstract:

This paper explores the role of artificial intelligence (AI) in enhancing TB detection through chest radiography, particularly in low-resource settings where diagnostic expertise is limited. The study underscores AI's potential to address gaps in healthcare delivery and improve early TB detection.

Introduction:

The authors discuss the global TB burden and emphasize the need for scalable diagnostic tools to combat the disease effectively. They highlight the inadequacy of current manual interpretation methods due to the scarcity of skilled radiologists in underserved regions. AI-based systems, particularly those leveraging deep learning, are proposed as a solution to streamline diagnostics and improve accessibility.

Background:

Chest X-rays (CXRs) are one of the most commonly used tools for TB detection. However, interpretation often requires specialized expertise, which is not always available in low-resource settings. AI-driven computer-aided diagnosis (CAD) systems can analyze CXRs efficiently, providing standardized and accurate results.

Methodology:

The study reviews various AI models designed for TB detection, focusing on their datasets, model architectures, and validation processes. Key aspects include:

- The use of convolutional neural networks (CNNs) for feature extraction and classification.
- Data augmentation techniques to improve model robustness.
- Evaluation metrics such as accuracy, sensitivity, and specificity.

Observation:

AI-based systems demonstrated high diagnostic accuracy and consistency, comparable to human radiologists. However, challenges such as data variability, model interpretability, and integration into clinical workflows were noted.

Results:

The findings suggest that AI integration in TB diagnostics can improve early detection rates, reduce diagnostic errors, and enhance accessibility in resource-constrained environments.

Paper 4: “Triage of Persons With Tuberculosis Symptoms Using Artificial Intelligence-Based Chest Radiograph Interpretation: A Cost-Effectiveness Analysis” by Nsengiyumva NP, et al.[4]**Abstract:**

This study evaluates the cost-effectiveness and diagnostic accuracy of AI-based systems for triaging individuals with TB symptoms in low-resource settings, comparing their performance to conventional methods.

Introduction:

The authors discuss the challenges of TB screening in resource-limited areas, such as the scarcity of trained radiologists and high costs. AI-based systems are proposed as a means to streamline triage processes and reduce healthcare costs.

Background:

Triage is a critical step in TB management, especially in high-burden settings. Manual interpretation of CXRs can be time-intensive and inconsistent, making automated systems an attractive alternative.

Methodology:

The study models the implementation of AI-based triage systems and compares their cost-effectiveness to traditional methods. Key features include:

- Simulation of AI-assisted workflows in various healthcare settings.
- Evaluation of cost savings, time efficiency, and detection rates.

Observation:

AI-based systems demonstrated faster processing times and higher diagnostic consistency compared to traditional methods. Cost-effectiveness was especially pronounced in high-volume settings.

Results:

The findings suggest that AI-based triage systems can reduce costs and improve diagnostic accuracy, particularly in resource-constrained environments.

Journal Article: “Synthetic Oversampling for Imbalanced Medical Datasets” by N.V. Chawla, et al.[5]

Abstract:

This article explores the Synthetic Minority Oversampling Technique (SMOTE), a method designed to address the challenges of class imbalance in machine learning, particularly in medical datasets. The focus is on how SMOTE enhances model performance by generating synthetic samples for underrepresented classes.

Introduction:

Class imbalance is a common issue in medical datasets, where minority classes, such as rare diseases, are often underrepresented. This imbalance leads to biased machine learning models, which may overlook critical cases. SMOTE is introduced as a solution that creates synthetic samples to improve class representation during training.

Background:

Traditional resampling methods, such as random oversampling or under sampling, have limitations, such as redundancy or loss of valuable data. SMOTE overcomes these issues by generating new samples based on interpolation between existing minority class instances, thereby preserving dataset diversity.

Methodology:

- SMOTE generates synthetic samples by interpolating between neighboring minority class instances in the feature space.
- The authors applied SMOTE to various imbalanced medical datasets, including chest X-rays, to evaluate its impact on classification models.

- Performance metrics such as recall, precision, and F1-score were used to assess the effectiveness of SMOTE.

Observation:

- SMOTE improved recall for minority classes significantly (up to 20%) without negatively affecting precision.
- The method was particularly effective for medical datasets with extreme class imbalances, such as TB-positive cases in chest X-rays.

Results:

The study validated SMOTE as an effective preprocessing technique for handling class imbalance in machine learning applications. It enables models to better detect minority class instances, which is critical in medical diagnostics.

Paper 6: "Deep Learning at Chest Radiography: Automated Classification of Pulmonary Tuberculosis by Using Convolutional Neural Networks" by Lakhani P, Sundaram B. [6]

Abstract:

This foundational study explores the efficacy of deep learning, specifically Convolutional Neural Networks (CNNs), for the automated classification of pulmonary tuberculosis (TB) from digital chest radiographs. The primary objective was to demonstrate the potential of these advanced computational methods to enhance diagnostic efficiency and accuracy, thereby addressing critical needs in global TB control, especially in resource-constrained environments where access to expert radiologists is limited. The study aimed to establish the feasibility of using AI to rapidly and accurately identify TB patterns on chest X-rays.

Introduction:

The authors begin by underscoring the severe global health challenge posed by tuberculosis, which continues to be a leading cause of morbidity and mortality worldwide. They highlight the significant difficulties inherent in its diagnosis, particularly in low- and middle-income countries where diagnostic infrastructure, including the availability of trained radiologists, is often insufficient. Manual interpretation of chest X-rays, a widely

used screening tool, is acknowledged as being subjective and time-consuming, prone to inter-reader variability. In this context, the paper proposes deep learning, and specifically CNNs, as a cutting-edge technological solution to automate and streamline the TB diagnostic process, offering the promise of rapid, consistent, and scalable interpretations that could revolutionize TB screening programs.

Background:

The background section elaborates on the existing diagnostic landscape for TB. Traditional methods of TB diagnosis from chest X-rays are characterized as labor-intensive and requiring specialized expertise, often leading to delays in diagnosis and treatment initiation. The inherent limitations of human visual interpretation, such as fatigue and varying levels of experience, make automated systems an attractive alternative. The authors frame their work within the broader context of the burgeoning field of medical image analysis using artificial intelligence. They specifically position CNNs as uniquely suited for image recognition tasks due to their ability to automatically learn hierarchical features from raw image data, obviating the need for manual feature engineering that often plagues traditional image processing techniques. This deep learning approach is presented as a paradigm shift towards achieving more consistent and efficient diagnostic outcomes.

Methodology:

The study employed a robust methodology for training and evaluating CNN models for TB classification. Key methodological features included:

- **Dataset Acquisition:** The researchers utilized a large dataset of de-identified digital chest radiographs, meticulously annotated for the presence or absence of active pulmonary tuberculosis. This dataset was carefully curated to ensure representativeness and diagnostic accuracy, often involving expert radiologist consensus.
- **CNN Architectures:** A variety of established and widely-used CNN architectures were explored and adapted for the task, including prominent models like AlexNet, GoogLeNet, and ResNet. The choice of diverse architectures allowed for a comprehensive evaluation of different network depths and complexities.
- **Training and Validation:** The dataset was rigorously split into training, validation,

and test sets to ensure unbiased evaluation of the models' generalization capabilities. Standard deep learning training protocols were followed, including backpropagation with stochastic gradient descent, appropriate loss functions (e.g., cross-entropy), and regularization techniques (e.g., dropout) to prevent overfitting.

- **Performance Metrics:** The performance of the trained CNN models was meticulously evaluated using a comprehensive suite of statistical metrics commonly employed in medical diagnostics. These included accuracy, sensitivity (recall), specificity, precision, F1-score, and Area Under the Receiver Operating Characteristic (ROC) curve (AUC). These metrics provided a nuanced understanding of the models' ability to correctly identify both positive (TB) and negative (non-TB) cases.
- **Comparative Analysis:** A crucial aspect of the methodology involved directly comparing the diagnostic performance of the CNNs against that of multiple experienced radiologists. This head-to-head comparison provided a benchmark for the AI system's capabilities relative to human experts.

Observation:

Through extensive experimentation and evaluation, the researchers observed compelling performance from the trained CNN models. The models consistently demonstrated promising capabilities in accurately distinguishing patterns indicative of active pulmonary TB from normal chest radiographs and those exhibiting other non-TB related abnormalities. A significant observation was that the diagnostic accuracy of the CNNs was not only comparable to but, in several instances, was found to marginally exceed that of individual human readers. This suggested a level of consistency and diagnostic power that could significantly enhance the efficiency of TB screening programs, particularly in scenarios where expert human resources are scarce. The models also exhibited robust performance across various subsets of the data, indicating good generalization.

Results:

The study concluded with strong evidence supporting the utility of deep learning, specifically CNNs, for the automated classification of pulmonary tuberculosis from chest radiographs. The findings definitively suggested that these AI-driven systems could effectively automate a critical step in the TB diagnostic pathway. The observed high accuracy, sensitivity, and specificity of the CNNs indicated their potential as a powerful

tool to not only augment the capabilities of human radiologists but also to potentially serve as a first-line screening mechanism. The authors posited that such automated systems could significantly accelerate TB diagnosis, reduce reliance on limited expert resources, and ultimately contribute to improved patient outcomes by facilitating earlier detection and treatment initiation, particularly in environments facing significant healthcare infrastructure challenges.

Paper 7: "Efficient deep network architectures for fast chest X-ray tuberculosis screening and visualization" by Pasa, F., Golkov, V., Pfeiffer, F., Cremers, D., & Pfeiffer, D. [2]

Abstract:

This research focuses on addressing two critical limitations of existing deep learning models for medical image analysis: computational efficiency and interpretability. The study aims to develop and rigorously evaluate novel, efficient deep network architectures specifically tailored for rapid and accurate chest X-ray tuberculosis (TB) screening. Crucially, it also integrates methods for visualizing the internal decision-making processes of these networks, thereby moving beyond the "black-box" nature often associated with deep learning models and enhancing their trustworthiness for clinical application.

Introduction:

The authors initiate their discussion by acknowledging the ongoing global public health crisis of tuberculosis and the pressing need for more scalable, faster, and universally accessible screening methods, particularly in high-burden regions where diagnostic resources are often overwhelmed. While recognizing the significant progress made by deep learning in medical image analysis, they identify a gap: many state-of-the-art models are computationally intensive, making their deployment in resource-limited settings challenging. Furthermore, the lack of transparency in how these models arrive at their conclusions—the "black-box" problem—is a significant barrier to clinical adoption. This paper thus proposes a two-pronged approach: designing computationally efficient deep neural networks for rapid deployment and integrating visualization techniques to illuminate the networks' decision-making processes, thereby fostering greater confidence and

interpretability.

Background:

The background section delves into the practical challenges of deploying deep learning in real-world clinical settings. While deep learning has demonstrated remarkable accuracy in tasks like medical image classification, the large model sizes and high computational demands of many leading architectures can hinder their implementation on less powerful hardware often found in clinics or mobile screening units. Moreover, the lack of transparency in deep learning models—where it's difficult to understand why a model made a particular prediction—is a major impediment to their acceptance by clinicians, who require a clear rationale for diagnostic decisions. The authors position their work as a vital step towards overcoming these hurdles by focusing on architecturally efficient models and providing interpretable insights into their predictions, thus bridging the gap between cutting-edge AI research and practical clinical utility.

Methodology:

The study employed a sophisticated methodology centered on the development and rigorous evaluation of specialized deep network architectures. Key methodological features included:

- **Efficient Network Design:** The core of the methodology involved designing and optimizing novel deep learning architectures with a focus on computational efficiency. This included exploring techniques such as depth-wise separable convolutions, neural architecture search, pruning, and quantization to reduce model parameters and computational complexity without significantly compromising accuracy. The aim was to create "lightweight" models suitable for fast inference on diverse hardware.
- **Large-Scale Dataset Utilization:** The models were trained and validated on extensive datasets of chest X-rays, ensuring a robust learning process and the ability to generalize across varied patient populations and image characteristics. Data augmentation techniques were employed to enhance the training set's diversity and prevent overfitting.

- **Visualization Techniques Integration:** A crucial methodological component was the integration of interpretable AI (XAI) techniques. This involved methods such as Gradient-weighted Class Activation Mapping (Grad-CAM), LIME (Local Interpretable Model-agnostic Explanations), or SHAP (SHapley Additive exPlanations). These techniques generate "saliency maps" or importance scores that highlight the specific regions in the input chest X-ray that most influenced the network's prediction for TB, providing a visual explanation for its decision.
- **Performance Evaluation:** Model performance was assessed using standard diagnostic metrics (accuracy, sensitivity, specificity, AUC) but also crucially evaluated for inference speed (e.g., images processed per second) and model size (e.g., number of parameters, memory footprint) to quantify their efficiency for real-time screening.
- **Comparative Analysis:** The performance of the newly designed efficient architectures was compared against established, larger deep learning models to demonstrate the trade-offs between efficiency and accuracy, aiming to achieve an optimal balance.

Observation:

The empirical observations from this study were highly encouraging. The researchers found that their newly developed and optimized efficient deep network architectures consistently achieved high accuracy in TB detection, comparable to or even exceeding that of much larger, more computationally expensive models. A particularly significant observation was the remarkable improvement in processing times, making these models highly suitable for rapid screening in high-throughput environments. Furthermore, the integrated visualization techniques proved highly effective. The generated saliency maps clearly highlighted clinically relevant areas on the chest X-rays (e.g., infiltrates, cavitations) that correlated with the network's TB prediction. This visual correlation provided compelling evidence that the networks were "looking at" and basing their decisions on medically meaningful features, significantly enhancing their transparency and interpretability for clinicians.

Results:

The findings conclusively demonstrated that it is indeed possible to develop efficient deep network architectures that enable both fast and accurate chest X-ray TB screening. The

study successfully showcased that reducing model complexity does not necessarily entail a significant drop in diagnostic performance, thereby making these AI systems practical for deployment even on less powerful computational hardware. Crucially, the integration of visualization methods was a significant breakthrough, effectively transforming these AI models from "black boxes" into more transparent and interpretable tools. By providing visual explanations for their predictions, these systems foster greater trust and facilitate clinical adoption. The authors concluded that these advancements represent a significant step towards enabling widespread, rapid, and trustworthy AI-powered TB screening solutions, which hold immense potential for impacting global health by accelerating diagnosis and facilitating timely intervention in high-burden settings.

2.1 Inference from Literature Survey

The collective insights from the reviewed studies clearly emphasize the growing relevance of artificial intelligence in medical diagnostics, particularly for tuberculosis screening using chest radiographs. These research efforts showcase that deep learning models—when properly trained, fine-tuned, and validated—can significantly improve the speed, accuracy, and reliability of TB diagnosis, even in resource-constrained healthcare settings.

Several core takeaways emerged:

- There is widespread acknowledgment of the limitations in traditional TB screening methods, such as reliance on radiologist expertise, time-consuming procedures, and variability in interpretations. This opens a clear need for automated and objective diagnostic systems.
- Across the surveyed works, **transfer learning, data augmentation, and ensemble modeling** have consistently been effective in improving classification performance, often exceeding 90% accuracy. These findings underline the **importance of model reusability and robust training pipelines**.
- Researchers also highlighted the need for practical deployment and real-world application, stressing that models must be lightweight, interpretable, and easy to integrate into frontline medical workflows—especially in TB-prone regions.

Thus, the literature collectively supports the core idea that intelligent diagnostic systems using CNNs and modern machine learning practices are no longer experimental—they are essential. The progress observed in prior work directly motivates the current research, which aims to develop a practical, tunable, and deployable TB detection system. By combining proven deep learning strategies with focused performance optimization and deployment tools (such as Streamlit), this project aspires to bridge the gap between academic models and real-world healthcare impact.

Chapter 3

3. Proposed Work

3.1 Introduction

This chapter outlines the proposed approach for developing a deep learning system for Tuberculosis detection from chest X-ray images, based on the problem statement and objectives defined in Chapter 1 and informed by the literature review in Chapter 2. The proposal details the planned steps to address the challenge of automated diagnosis, starting from the raw data format and progressing through model development, evaluation, and deployment.

3.2 Proposed Work

The proposed work involved implementing a multi-stage pipeline designed to build and evaluate a high-performing deep learning model for binary classification of chest X-rays (Normal vs. Tuberculosis). The key components of the proposal were:

- **Raw Data Preprocessing & Conversion:** It was proposed that raw medical images, initially in the DICOM format, would be converted into a more accessible image format like PNG. This step was necessary to interface with standard image processing and deep learning libraries. The conversion process would need to handle medical-specific aspects such as applying rescale slope and intercept and ensuring correct pixel value normalization. The converted data would be organized into class-specific directories.
- **Dataset Structuring:** The proposed methodology included structuring the dataset into distinct and non-overlapping subsets for training, validation, and final testing. A standard split ratio (e.g., 80% train, 10% validation, 10% test) would be applied to allow for proper

model development and evaluation workflow, ensuring the test set remained truly unseen during training and tuning.

- **Model Selection through Comparative Evaluation:** Capitalizing on the power of transfer learning, it was proposed to evaluate multiple state-of-the-art Convolutional Neural Network (CNN) architectures pre-trained on large datasets like ImageNet. These models would be used as fixed feature extractors by replacing their final classification layers with a new head suitable for binary classification. A robust evaluation strategy, specifically Stratified K-Fold Cross-Validation, was proposed to compare the performance of these different architectures across the *entire* available dataset, providing a more reliable performance estimate than a single train/test split. This comparative analysis would aim to identify the most promising architecture(s) for the task.
- **Hyperparameter Optimization:** Following the selection of a promising model architecture from the comparative evaluation, it was proposed to optimize the hyperparameters of the newly added classification layers and the training process (e.g., learning rate, dropout rates, number of dense units). An automated hyperparameter tuning method, such as Hyperband, was proposed to efficiently search the parameter space using the designated training and validation sets, with the objective of maximizing a key performance metric (e.g., validation AUC).
- **Final Model Training:** The proposal included training the final version of the selected and tuned model. This final training would be performed on a larger dataset, typically the combination of the training and validation sets used during tuning, using the optimal hyperparameters found in the previous step. This step aims to leverage as much data as possible for the final model.
- **Rigorous Performance Evaluation:** After training the final model, a critical step in the proposal was to evaluate its performance comprehensively on the strictly held-out test set. This evaluation would use a standard set of classification metrics, including Accuracy, F1 Score, ROC AUC, PR AUC, and analysis of the Confusion Matrix, to provide an unbiased assessment of the model's expected performance in a real-world scenario.

- **Model Deployment:** Finally, it was proposed to deploy the trained model in a basic, interactive format to demonstrate its practical utility. This could involve creating a simple web application where users could upload images and receive predictions from the model. This proposed work outlined the planned steps and techniques intended to achieve the project's objectives, forming the blueprint for the implementation described in the subsequent chapters.

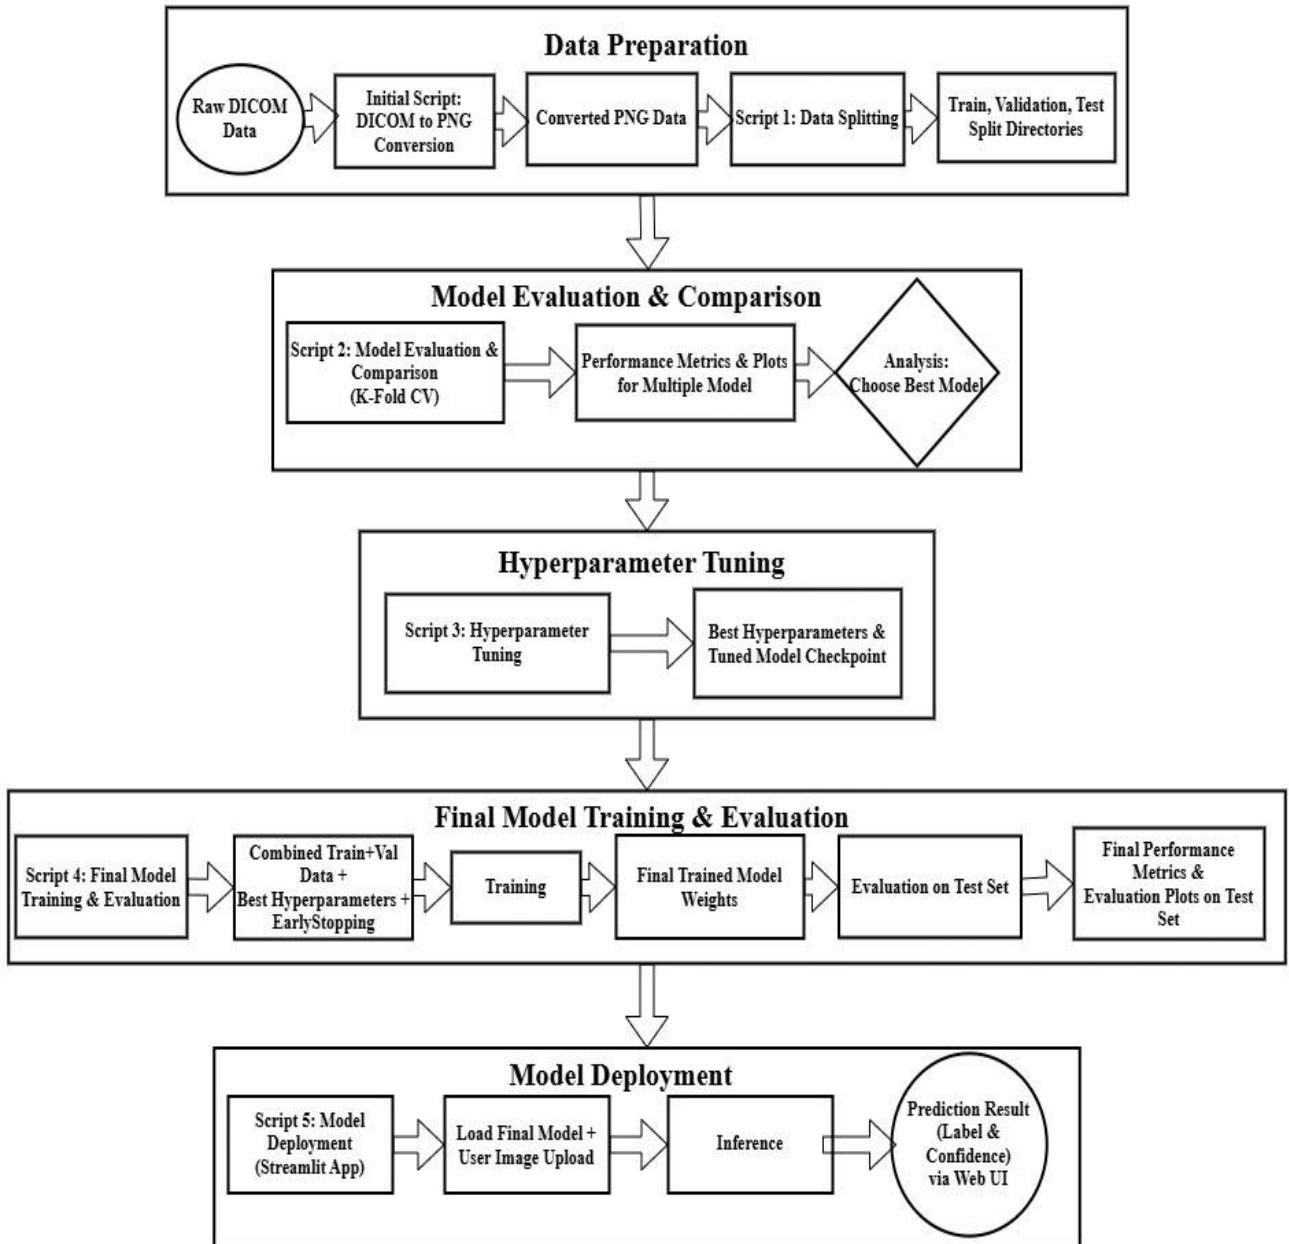


Fig 3.1 Workflow Diagram of the Proposed TB Detection System

This flowchart details a comprehensive machine learning pipeline, emphasizing a structured approach from raw data ingestion to user-friendly model deployment.

The workflow is meticulously divided into five key phases:

Data Preparation:

- **Initial Data Ingestion:** The process commences with raw DICOM images which are commonly used medical images containing rich metadata.
- **Format Transformation:** An initial Script: DICOM to PNG Conversion processes these raw DICOM files, converting them into PNG format. PNG is a lossless image format, suitable for machine learning input. This step ensures data compatibility with subsequent processing stages.
- **Dataset Stratification:** Script 1: Data Splitting script meticulously partitions the converted PNG data into three distinct sets: "Train, Validation, Test Split Directories." This critical step is vital for preventing data leakage and ensuring an unbiased evaluation of the model's generalization capabilities. The training set is used for model learning, the validation set for hyperparameter tuning and early stopping, and the test set for final, unseen performance assessment.

Model Evaluation & Comparison:

- **Multi-Model Assessment and Cross-Validation:** "Script 2: Model Evaluation & Comparison (K-Fold CV)" executes the core evaluation. This script rigorously assesses the performance of multiple candidate machine learning models or different model architectures using K-Fold Cross-Validation. K-Fold CV is a robust resampling technique that helps estimate the model's true performance by training and validating it on different subsets of the training data, reducing variance.
- **Performance Visualization and Metrics:** The execution generates various Performance Metrics & Plots for Multiple Models. These include quantitative measures (accuracy, precision, recall, F1-score) and visual representations (ROC curves, confusion matrices) that allow for a comparative analysis of the models.
- **Optimal Model Selection:** An Analysis phase critically reviews these metrics and plots to choose the best model. This decision is driven by the defined project objectives and the performance observed across various metrics.

Hyperparameter Tuning:

- **Optimization Script Execution:** "Script 3: Hyperparameter Tuning" is invoked, focusing on the model selected in the previous stage. This script systematically explores various combinations of the chosen model's hyperparameters (parameters external to the model, such as learning rate, batch size, number of layers, etc., which are set before the training process).
- **Generation of Optimized Model:** The outcome of this tuning process is Best Hyperparameters & Tuned Model Checkpoint. This checkpoint represents the model's configuration with the most effective hyperparameters identified, ready for final training.

Final Model Training & Evaluation:

- **Comprehensive Training Phase:** "Script 4: Final Model Training & Evaluation" initiates the definitive training of the model.
- **Leveraging Combined Data and Optimized Settings:** The model is trained on combined Train+Val Data, utilizing a larger dataset for learning, and importantly, incorporates the best hyperparameters derived from the previous tuning phase.
- **Regularization via Early Stopping:** "EarlyStopping" is strategically implemented during this training. This is a regularization technique that monitors the model's performance on the validation set and automatically halts training when improvement on this set ceases, preventing overfitting and ensuring the model generalizes well to new data.
- **Persisting Trained Model:** The result of this training is Final Trained Model Weights, which encapsulates the learned parameters of the optimized model.
- **Unbiased Performance Assessment:** A crucial evaluation on Test Set is then performed. This step uses the completely unseen and untouched test data to generate final Performance Metrics & evaluation plots on test set. This provides an unbiased and definitive assessment of the model's real-world predictive accuracy and robustness.

Model Deployment:

- **Application Development:** "Script 5: Model Deployment (Streamlit App)"

facilitates the operationalization of the trained model. Streamlit is a Python library often used to create interactive web applications for data science and machine learning.

- **Model Integration and User Input:** This stage involves integrating final model into the Streamlit application and enabling user image upload, allowing end-users to submit new images for prediction.
- **Real-time Prediction:** The deployed model then performs inference on the user-uploaded images, applying its learned knowledge to generate predictions.
- **User-Friendly Output:** The ultimate outcome is a Prediction Result (Label & Confidence) via Web UI. This provides users with the model's prediction (a classification label for medical diagnosis) along with a Confidence score, indicating the model's certainty, all delivered through an intuitive web-based interface.

Chapter 4

4. Methodology

4.1 Introduction

This chapter details the comprehensive methodology employed in this project to implement the proposed work outlined in Chapter 3. It describes how each stage of the proposed pipeline, from raw data handling to model deployment, was executed using a set of six distinct scripts. The methodology provides a step-by-step account of the processes, techniques, and tools used to develop, evaluate, and deploy the deep learning model for Tuberculosis detection.

4.2 Implementation Strategy

The implementation strategy followed the proposed multi-stage pipeline sequentially, with each stage implemented by one or more of the developed scripts. The outputs from earlier stages served as necessary inputs for subsequent stages, creating a directed workflow from raw data to a deployable model. The implementation was executed through the following stages given in the flowchart:

4.2.1 Flowchart of the Proposed System

The below flowchart presents the complete workflow adopted in this project. Each step has been expanded in the subsequent sections below.

It describes how each stage of the proposed pipeline, from raw data handling and preparation to advanced deep learning model development, evaluation, and eventual deployment, was systematically executed. A series of six distinct scripts were developed and utilized, each serving a specific function in this multi-stage approach. This chapter provides a step-by-step account of the processes, techniques, and tools used to realize the project's objectives in building a deep learning model for Tuberculosis detection from chest radiography images.

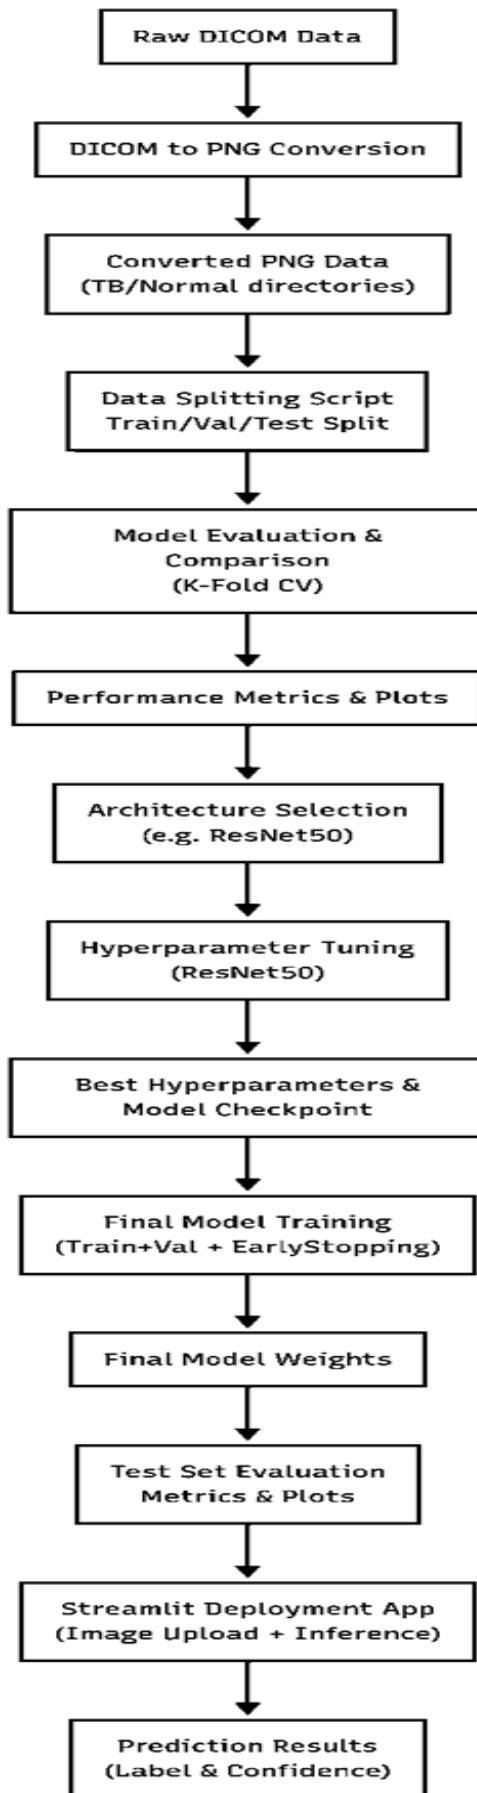


Fig 4.1 – Flowchart of the TB Detection Methodology Pipeline

4.3 Methodology Steps

4.3.1. Raw Data Handling and Conversion:

- **Objective:** To implement the proposed conversion of raw medical images from the DICOM format into PNG and organize them by class, preparing the data for standard image processing and deep learning workflows.
- **Script(s) Involved:** Initial DICOM Conversion Script (di_convert.py)

```
def convert_dicoms_to_png(source_folder, target_tb_folder, prefix="TB-"):  
    os.makedirs(target_tb_folder, exist_ok=True)  
    index = get_next_index_tb(target_tb_folder, prefix=prefix)  
  
    for filename in sorted(os.listdir(source_folder)):  
        if filename.lower().endswith((".dcm", ".dicom")):  
            dicom_path = os.path.join(source_folder, filename)  
            output_filename = f"{prefix}{index:05d}.png"  
            output_path = os.path.join(target_tb_folder, output_filename)  
  
            try:  
                ds = pydicom.dcmread(dicom_path)  
                pixel_array = ds.pixel_array.astype(np.float32)  
  
                # Apply rescale slope/intercept if present  
                intercept = ds.get("RescaleIntercept", 0.0)  
                slope = ds.get("RescaleSlope", 1.0)  
                pixel_array = pixel_array * slope + intercept
```

Fig 4.2 – Code Snippet of DICOM Conversion

- **Input:** Source folders containing chest X-ray images in DICOM (.dcm, .dicom) format.

📁	intbtr29	5/1/2025 12:01 PM
📁	intbtr65	5/1/2025 12:01 PM
📁	intbtr67	5/1/2025 3:42 PM
📁	intbtr78	5/1/2025 12:01 PM
📁	intbtr79	5/1/2025 12:01 PM
📁	intbtr81	5/1/2025 12:01 PM
📁	intbtr103	5/1/2025 12:01 PM
📁	intbtr104	5/1/2025 12:01 PM
📁	intbtr251	5/1/2025 12:01 PM
📁	intbtr252	5/1/2025 12:01 PM
📁	intbtr253	5/1/2025 12:01 PM

Fig 4.3 Snippet of Source Folders

- **Process:**

- The Initial DICOM Conversion Script was implemented to iterate through predefined source directories containing DICOM files for both TB and Normal cases.
- For each DICOM file, pydicom was used to read the data, extract pixel arrays, and apply the **RescaleIntercept** and **RescaleSlope** to correctly scale pixel intensities.
- Pixel data was normalized to the 0-255 range, converted to **uint8**, and transformed to RGB format if originally gray scale.
- The processed images were saved as PNG files into designated target directories (e.g., **TB_Chest_Radiography_Database_raw/TB**, **TB_Chest_Radiography_Database_raw/Normal**) using the Pillow library, with sequential naming.
- **Output:** A directory structure (**TB_Chest_Radiography_Database_raw**) populated with converted PNG images, organized into TB and Normal subdirectories.

```
Saved TB-06324.png
Saved TB-06325.png
Saved TB-06326.png
Saved TB-06327.png
Saved TB-06328.png
Saved TB-06329.png
Saved TB-06330.png
Saved TB-06331.png
Saved TB-06332.png
TB Conversion complete!
Starting Normal dicom files conversion to PNG...
🔍 Processing folder: intbtr12
Saved Normal-00001.png
Saved Normal-00002.png
Saved Normal-00003.png
Saved Normal-00004.png
Saved Normal-00005.png
Saved Normal-00006.png
```

Fig 4.4 – Snippet of DICOM Conversion Terminal

4.3.2. Structured Dataset Creation:

- **Objective:** To implement the proposed partitioning of the converted raw image dataset into training, validation, and test sets.
- **Script(s) Involved:** Script for Data Splitting (data_split.py).

```
# --- Split Dataset (Run Once) ---
def split_dataset():
    if os.path.exists(OUTPUT_DIR):
        print("Dataset already split.")
        return

    print("Splitting dataset into train/val/test...")
    os.makedirs(OUTPUT_DIR, exist_ok=True)
    for cls in ["Normal", "Tuberculosis"]:
        files = os.listdir(os.path.join(ORIGINAL_DIR, cls))
        random.seed(SEED)
        random.shuffle(files)

        total = len(files)
        train_end = int(total * SPLIT_RATIO[0])
        val_end = train_end + int(total * SPLIT_RATIO[1])

        split_files = {
            "train": files[:train_end],
            "val": files[train_end:val_end],
            "test": files[val_end:]
        }

        for split in split_files:
            split_dir = os.path.join(OUTPUT_DIR, split, cls)
            os.makedirs(split_dir, exist_ok=True)
            for file in tqdm(split_files[split], desc=f"{cls} - {split}"):
                shutil.copy2(os.path.join(ORIGINAL_DIR, cls, file), os.path.join(split_dir, file))
```

Fig 4.5 – Code Snippet of Data Splitting

- **Input:** The raw image dataset directory (**TB_Chest_Radiography_Database_raw**) containing PNG images organized by class, as created in section 4.3.1
- **Process:**
 - Data_split.py was implemented to read filenames from the raw class directories.
 - Image filenames for each class were shuffled using a fixed random seed for reproducibility.
 - The shuffled lists were split into train, validation, and test sets based on predefined ratios (80/10/10).

- A new output directory structure (**TB_Chest_Radiography_Database_split**) with **train**, **val**, **test** subdirectories, each containing class subdirectories, was created.
- Files were copied from the raw directories to their corresponding locations in the new split structure.
- **Output: A structured dataset directory (TB_Chest_Radiography_Database_split)** containing distinct **train**, **val**, and **test** sets, ready for model input.

```
PS C:\Users\shivi\OneDrive\Desktop\tb_mod_comp> & c:/Users/shivi/OneDrive/Desktop/tb_mod_comp/cmpvenv/python.exe c:/Users/shivi/OneDrive/Desktop/tb_mod_comp/data_split.py
Splitting dataset into train/val/test...
Normal - train: 100%|██████████| 7627/7627 [00:32<00:00, 231.41it/s]
Normal - val: 100%|██████████| 953/953 [00:04<00:00, 231.69it/s]
Normal - test: 100%|██████████| 954/954 [00:04<00:00, 228.73it/s]
Tuberculosis - train: 100%|██████████| 7628/7628 [00:31<00:00, 241.74it/s]
Tuberculosis - val: 100%|██████████| 952/952 [00:03<00:00, 245.76it/s]
Tuberculosis - test: 100%|██████████| 954/954 [00:03<00:00, 239.87it/s]
PS C:\Users\shivi\OneDrive\Desktop\tb_mod_comp>
```

Fig 4.6 – Snippet of Data Splitting terminal

4.3.3. Exploratory Model Evaluation and Selection:

- **Objective:** To implement the proposed comparative evaluation of multiple CNN architectures using transfer learning and K-Fold Cross-Validation to select promising models.
- **Script(s) Involved:** Script for Model Evaluation and Comparison (model_comp).
- **Input:** The structured dataset directory (**TB_Chest_Radiography_Database_split**) created in section 4.3.2

- **Process:**

- Script was implemented to load all image paths and labels from the split dataset.
- Functions were defined to build transfer learning models using Keras Applications (ResNet50, Xception, etc.), freezing base weights and adding a custom classification head.
- **tf.data.Dataset** pipelines were configured for data loading and model-specific preprocessing.

```

def train_on_fold(X, y, model_name):
    model_fn, preprocess_fn = get_model_fn(model_name)
    skf = StratifiedKFold(n_splits=FOLDS, shuffle=True, random_state=SEED)
    metrics = []
    model_times = []
    all_y_true, all_y_pred = [], []
    model_conf_matrices = []

    for fold, (train_idx, test_idx) in enumerate(skf.split(X, y)):
        print(f"\nFold {fold+1}/{FOLDS} - {model_name}")

        plot_class_distribution(y[train_idx], f"{model_name} Fold {fold+1} - Train Class Distribution", f"{model_name}_fold{fold+1}_train_distribution.png")
        plot_class_distribution(y[test_idx], f"{model_name} Fold {fold+1} - Test Class Distribution", f"{model_name}_fold{fold+1}_test_distribution.png")

        X_train, y_train = X[train_idx], y[train_idx]
        X_test, y_test = X[test_idx], y[test_idx]

        train_ds = create_dataset(X_train, y_train, preprocess_fn, shuffle=True)
        test_ds = create_dataset(X_test, y_test, preprocess_fn, shuffle=False)

        count_class_distribution(train_ds, "Training")
        count_class_distribution(test_ds, "Testing")

        model = build_model(model_fn)
        model.compile(optimizer="adam", loss="binary_crossentropy", metrics=["accuracy", AUC(name="auc")])

        checkpoint_path = f"model_{model_name}_fold{fold+1}.h5"
        checkpoint = ModelCheckpoint(checkpoint_path, monitor='val_auc', save_best_only=True, mode='max', verbose=1)
        early_stop = EarlyStopping(patience=3, restore_best_weights=True, monitor='val_loss', mode='min', verbose=1)

        start_time = time.time()
        history = model.fit(train_ds, validation_data=test_ds, epochs=EPOCHS, callbacks=[early_stop, checkpoint], verbose=1)
        elapsed_time = time.time() - start_time
        model_times.append(elapsed_time)

```

Fig 4.7 – Code Snippet of K-Fold Cross-Validation Function

```

def create_dataset(paths, labels, preprocess_fn, shuffle=False):
    def _load_image(path, label):
        image = tf.io.read_file(path)
        image = tf.image.decode_image(image, channels=3)
        image.set_shape([None, None, 3])
        image = tf.image.resize(image, IMG_SIZE)
        image = tf.cast(image, tf.float32)
        image = preprocess_fn(image)
        return image, label

    ds = tf.data.Dataset.from_tensor_slices((paths, labels))
    if shuffle:
        ds = ds.shuffle(buffer_size=len(paths), seed=SEED)
    ds = ds.map(_load_image, num_parallel_calls=tf.data.AUTOTUNE)
    return ds.batch(BATCH_SIZE).prefetch(tf.data.AUTOTUNE)

```

Fig 4.8 – Code Snippet of Dataset tensor creation

- Stratified K-Fold cross-validation (5 folds) was performed on the entire dataset, ensuring class distribution in folds.
- Each predefined model architecture was trained and evaluated iteratively over the 5 folds. Training used fold-specific train data and validated on fold-specific test data, with **EarlyStopping** and **ModelCheckpoint**.
- Performance metrics (Accuracy, F1, ROC AUC, PR AUC), Confusion Matrices, and Classification Reports were calculated and collected per fold.
- Memory was managed between folds and models.
- Mean and standard deviation of metrics across folds were computed for each model.
- Results were compiled, saved to CSVs, and comparative plots were generated.
- The script identified the best performing models based on aggregated metrics, informing the choice for the next stage.
- **Output: Detailed performance metrics across K-Fold CV for multiple models, comparative plots, CSV summaries, and identification of promising model architectures (ResNet50).**

```

Epoch 16: val_auc did not improve from 0.96057
1906/1906 [=====] - 82s 43ms/step - loss: 0.2988 - accuracy: 0.8700 - auc: 0.9452 - val_loss: 0.2600 - val_accuracy: 0.8917 - val_auc: 0.9590
Epoch 17/18
1906/1906 [=====] - ETA: 0s - loss: 0.2969 - accuracy: 0.8708 - auc: 0.9457Restoring model weights from the end of the best epoch: 14.

Epoch 17: val_auc did not improve from 0.96057
1906/1906 [=====] - 83s 43ms/step - loss: 0.2969 - accuracy: 0.8708 - auc: 0.9457 - val_loss: 0.2575 - val_accuracy: 0.8948 - val_auc: 0.9595
Epoch 17: early stopping
477/477 [=====] - 18s 36ms/step
Memory cleared for EfficientNetB0 after fold 2...

Fold 3/5 - EfficientNetB0
c:\Users\shivi\OneDrive\Desktop\tb_mod_comp\model_comp.py:120: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

sns.barplot(x=list(class_counts.keys()), y=list(class_counts.values()), palette="pastel")
c:\Users\shivi\OneDrive\Desktop\tb_mod_comp\model_comp.py:120: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

sns.barplot(x=list(class_counts.keys()), y=list(class_counts.values()), palette="pastel")
Training split class counts: Normal=7627, TB=7621
Testing split class counts: Normal=1907, TB=1905
Epoch 1/18
1906/1906 [=====] - ETA: 0s - loss: 0.3935 - accuracy: 0.8211 - auc: 0.9036
Epoch 1: val_auc improved from -inf to 0.95240, saving model to model_EfficientNetB0_fold3.h5
1906/1906 [=====] - 90s 45ms/step - loss: 0.3935 - accuracy: 0.8211 - auc: 0.9036 - val_loss: 0.2924 - val_accuracy: 0.8720 - val_auc: 0.9524
Epoch 2/18

```

Fig 4.9 – Code Snippet for Fold 3 of EfficientNetB0 training

```

Fold 2/5 - MobileNetV2
c:\Users\shivi\OneDrive\Desktop\tb_mod_comp\model_comp.py:120: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

sns.barplot(x=list(class_counts.keys()), y=list(class_counts.values()), palette="pastel")
c:\Users\shivi\OneDrive\Desktop\tb_mod_comp\model_comp.py:120: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

sns.barplot(x=list(class_counts.keys()), y=list(class_counts.values()), palette="pastel")
Training split class counts: Normal=7627, TB=7621
Testing split class counts: Normal=1907, TB=1905
Epoch 1/18
1906/1906 [=====] - ETA: 0s - loss: 0.3656 - accuracy: 0.8297 - auc: 0.9156
Epoch 1: val_auc improved from -inf to 0.95330, saving model to model_MobileNetV2_fold2.h5
1906/1906 [=====] - 61s 31ms/step - loss: 0.3656 - accuracy: 0.8297 - auc: 0.9156 - val_loss: 0.2785 - val_accuracy: 0.8846 - val_auc: 0.9533
Epoch 2/18
1905/1906 [=====>.] - ETA: 0s - loss: 0.3082 - accuracy: 0.8657 - auc: 0.9413
Epoch 2: val_auc improved from 0.95330 to 0.95768, saving model to model_MobileNetV2_fold2.h5
1906/1906 [=====] - 58s 30ms/step - loss: 0.3082 - accuracy: 0.8657 - auc: 0.9413 - val_loss: 0.2697 - val_accuracy: 0.8856 - val_auc: 0.9577
Epoch 3/18

```

Fig 4.10 – Code Snippet for Fold 2 of MobileNetV2 training

```

Fold 4/5 - ResNet50
c:\Users\shivi\OneDrive\Desktop\tb_mod_comp\model_comp.py:120: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

sns.barplot(x=list(class_counts.keys()), y=list(class_counts.values()), palette="pastel")
c:\Users\shivi\OneDrive\Desktop\tb_mod_comp\model_comp.py:120: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

sns.barplot(x=list(class_counts.keys()), y=list(class_counts.values()), palette="pastel")
Training split class counts: Normal=7627, TB=7621
Testing split class counts: Normal=1907, TB=1905
Epoch 1/18
1906/1906 [=====] - ETA: 0s - loss: 0.3867 - accuracy: 0.8249 - auc: 0.9094
Epoch 1: val_auc improved from -inf to 0.95450, saving model to model_ResNet50_fold4.h5
1906/1906 [=====] - 90s 46ms/step - loss: 0.3867 - accuracy: 0.8249 - auc: 0.9094 - val_loss: 0.2815 - val_accuracy: 0.8822 - val_auc: 0.9545
Epoch 2/18
1905/1906 [=====>.] - ETA: 0s - loss: 0.3136 - accuracy: 0.8648 - auc: 0.9413
Epoch 2: val_auc improved from 0.95450 to 0.95766, saving model to model_ResNet50_fold4.h5
1906/1906 [=====] - 87s 46ms/step - loss: 0.3136 - accuracy: 0.8648 - auc: 0.9413 - val_loss: 0.2633 - val_accuracy: 0.8893 - val_auc: 0.9577
Epoch 3/18
1905/1906 [=====>.] - ETA: 0s - loss: 0.3088 - accuracy: 0.8701 - auc: 0.9439
Epoch 3: val_auc improved from 0.95766 to 0.96039, saving model to model_ResNet50_fold4.h5
1906/1906 [=====] - 88s 46ms/step - loss: 0.3089 - accuracy: 0.8700 - auc: 0.9439 - val_loss: 0.2977 - val_accuracy: 0.8743 - val_auc: 0.9604
Epoch 4/18
1906/1906 [=====] - ETA: 0s - loss: 0.3092 - accuracy: 0.8709 - auc: 0.9445
Epoch 4: val_auc improved from 0.96039 to 0.96212, saving model to model_ResNet50_fold4.h5
1906/1906 [=====] - 88s 46ms/step - loss: 0.3092 - accuracy: 0.8709 - auc: 0.9445 - val_loss: 0.2497 - val_accuracy: 0.8948 - val_auc: 0.9621
Epoch 5/18

```

Fig 4.11 – Code Snippet for Fold 4 of Resnet50 training

```

Epoch 15: val_auc did not improve from 0.94518
1906/1906 [=====] - 112s 59ms/step - loss: 0.3736 - accuracy: 0.8332 - auc: 0.9139 - val_loss: 0.3324 - val_accuracy: 0.8544 - val_auc: 0.9449
Epoch 16/18
1905/1906 [=====>.] - ETA: 0s - loss: 0.3697 - accuracy: 0.8352 - auc: 0.9153Restoring model weights from the end of the best epoch: 13.

Epoch 16: val_auc did not improve from 0.94518
1906/1906 [=====] - 112s 58ms/step - loss: 0.3696 - accuracy: 0.8353 - auc: 0.9154 - val_loss: 0.3163 - val_accuracy: 0.8583 - val_auc: 0.9445
Epoch 16: early stopping
477/477 [=====] - 22s 45ms/step
Memory cleared for DenseNet121 after fold 2...

💡 Fold 3/5 - DenseNet121
c:\Users\shivi\OneDrive\Desktop\tb_mod_comp\model_comp.py:120: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

sns.barplot(x=list(class_counts.keys()), y=list(class_counts.values()), palette="pastel")
c:\Users\shivi\OneDrive\Desktop\tb_mod_comp\model_comp.py:120: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

sns.barplot(x=list(class_counts.keys()), y=list(class_counts.values()), palette="pastel")
Training split class counts: Normal=7627, TB=7621
Testing split class counts: Normal=1907, TB=1905
Epoch 1/18
1905/1906 [=====>.] - ETA: 0s - loss: 0.4659 - accuracy: 0.7709 - auc: 0.8592
Epoch 1: val_auc improved from -inf to 0.94061, saving model to model_DenseNet121_fold3.h5
1906/1906 [=====] - 114s 58ms/step - loss: 0.4661 - accuracy: 0.7709 - auc: 0.8590 - val_loss: 0.3276 - val_accuracy: 0.8639 - val_auc: 0.9406
Epoch 2/18
1905/1906 [=====>.] - ETA: 0s - loss: 0.4000 - accuracy: 0.8178 - auc: 0.8999
Epoch 2: val_auc improved from 0.94061 to 0.94307, saving model to model_DenseNet121_fold3.h5
1906/1906 [=====] - 111s 58ms/step - loss: 0.4000 - accuracy: 0.8177 - auc: 0.8999 - val_loss: 0.3305 - val_accuracy: 0.8520 - val_auc: 0.9431

```

Fig 4.12 – Code Snippet for Fold 3 of DenseNet121 training

```

Epoch 4/18
1905/1906 [=====>.] - ETA: 0s - loss: 0.4176 - accuracy: 0.8144 - auc: 0.8955
Epoch 4: val_auc did not improve from 0.93041
1906/1906 [=====] - 110s 58ms/step - loss: 0.4175 - accuracy: 0.8145 - auc: 0.8956 - val_loss: 0.4065 - val_accuracy: 0.8161 - val_auc: 0.9268
Epoch 5/48
1905/1906 [=====>.] - ETA: 0s - loss: 0.4059 - accuracy: 0.8219 - auc: 0.9016Restoring model weights from the end of the best epoch: 2.

Epoch 5: val_auc did not improve from 0.93041
1906/1906 [=====] - 110s 58ms/step - loss: 0.4058 - accuracy: 0.8219 - auc: 0.9017 - val_loss: 0.3536 - val_accuracy: 0.8494 - val_auc: 0.9298
Epoch 5: early stopping
477/477 [=====] - 25s 50ms/step
Memory cleared for InceptionV3 after fold 2...

💡 Fold 3/5 - InceptionV3
c:\Users\shivi\OneDrive\Desktop\tb_mod_comp\model_comp.py:120: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

sns.barplot(x=list(class_counts.keys()), y=list(class_counts.values()), palette="pastel")
c:\Users\shivi\OneDrive\Desktop\tb_mod_comp\model_comp.py:120: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

sns.barplot(x=list(class_counts.keys()), y=list(class_counts.values()), palette="pastel")
Training split class counts: Normal=7627, TB=7621
Testing split class counts: Normal=1907, TB=1905
Epoch 1/18
1906/1906 [=====] - ETA: 0s - loss: 0.4972 - accuracy: 0.7583 - auc: 0.8422
Epoch 1: val_auc improved from -inf to 0.91993, saving model to model_InceptionV3_fold3.h5
1906/1906 [=====] - 96s 49ms/step - loss: 0.4972 - accuracy: 0.7583 - auc: 0.8422 - val_loss: 0.3616 - val_accuracy: 0.8402 - val_auc: 0.9199
Epoch 2/18
1906/1906 [=====] - ETA: 0s - loss: 0.4318 - accuracy: 0.8022 - auc: 0.8858

```

Fig 4.13 – Code Snippet for Fold 3 of InceptionV3 training

```

Epoch 13: val_auc did not improve from 0.94315
1906/1906 [=====] - 111s 58ms/step - loss: 0.3360 - accuracy: 0.8526 - auc: 0.9314 - val_loss: 0.3096 - val_accuracy: 0.8709 - val_auc: 0.9408
Epoch 13: early stopping
477/477 [=====] - 25s 51ms/step
Memory cleared for Xception after fold 2...

💡 Fold 3/5 - Xception
c:\Users\shivi\OneDrive\Desktop\tb_mod_comp\model_comp.py:120: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

sns.barplot(x=list(class_counts.keys()), y=list(class_counts.values()), palette="pastel")
c:\Users\shivi\OneDrive\Desktop\tb_mod_comp\model_comp.py:120: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

sns.barplot(x=list(class_counts.keys()), y=list(class_counts.values()), palette="pastel")
Training split class counts: Normal=7627, TB=7621
Testing split class counts: Normal=1907, TB=1905
Epoch 1/18
1905/1906 [=====>..] - ETA: 0s - loss: 0.4182 - accuracy: 0.8100 - auc: 0.8908
Epoch 1: val_auc improved from -inf to 0.93047, saving model to model_Xception_fold3.h5
1906/1906 [=====] - 111s 57ms/step - loss: 0.4182 - accuracy: 0.8100 - auc: 0.8908 - val_loss: 0.3423 - val_accuracy: 0.8473 - val_auc: 0.9305
Epoch 2/18

```

Fig 4.14 – Code Snippet for Fold 3 of Xception training

```

💡 Best model based on ROC-AUC + PR-AUC fallback: ResNet50

accuracy_mean accuracy_std f1_mean f1_std roc_auc_mean roc_auc_std pr_auc_mean pr_auc_std training_time_mean training_time_std pr_roc_delta
ResNet50 0.899528 0.002869 0.897901 0.003422 0.964326 0.001666 0.967322 0.001404 858.20079 162.659062 0.002995

The best model is: ResNet50
PS C:\Users\shivi\OneDrive\Desktop\tb_mod_comp>

```

Fig 4.15 – Snippet of Model Comparison terminal

4.3.4. Targeted Hyperparameter Optimization:

- **Objective:** To implement the proposed hyperparameter tuning for the custom classification head of the selected model architecture (ResNet50).
- **Script(s) Involved:** Script for Hyperparameter Tuning (hypertune_final.py).
- **Input:** The **train** and **val** splits from the structured dataset directory created in section 4.3.2, and the selected model architecture (ResNet50) identified in 4.3.3.
- **Process:**
 - hypertune_final.py was implemented to load the train and val datasets using **tf.keras.preprocessing.image_dataset_from_directory**.
 - tf.data.Dataset pipelines included ResNet-specific preprocessing.

```

# Build model with tunable hyperparameters
def build_model(hp):
    base_model = ResNet50(include_top=False, input_shape=IMG_SIZE + (3,), weights="imagenet")
    base_model.trainable = False

    x = layers.GlobalAveragePooling2D()(base_model.output)
    x = layers.Dropout(hp.Float('dropout', min_value=0.1, max_value=0.6, step=0.1))(x) # Expanded range
    x = layers.Dense(hp.Int('dense_units', min_value=128, max_value=1024, step=128), activation='relu')(x) # Expanded range
    x = layers.Dropout(hp.Float('dropout_2', min_value=0.1, max_value=0.6, step=0.1))(x) # Expanded range
    output = layers.Dense(1, activation='sigmoid')(x)

    model = Model(inputs=base_model.input, outputs=output)

    model.compile(
        optimizer=keras.optimizers.Adam(
            hp.Choice('learning_rate', [0.005, 0.003, 0.001, 0.0005]) # Adjusted learning rate values
        ),
        loss='binary_crossentropy',
        metrics=['accuracy', 'AUC', 'Precision', 'Recall'] # Added more metrics
    )

    return model

# Hyperband tuner for optimizing val_auc
tuner = Hyperband(
    build_model,
    objective=Objective("val_auc", direction="max"),
    max_epochs=30,
    factor=4,
    directory='resnet50_tuning_v2',
    project_name='tb_classification_v2'
)

```

Fig 4.16 – Code Snippet of Hyperband definition

```

# Load datasets
def load_datasets():
    train_ds = tf.keras.preprocessing.image_dataset_from_directory(
        os.path.join(OUTPUT_DIR, "train"), seed=SEED, image_size=IMG_SIZE,
        batch_size=BATCH_SIZE, label_mode="binary")
    val_ds = tf.keras.preprocessing.image_dataset_from_directory(
        os.path.join(OUTPUT_DIR, "val"), seed=SEED, image_size=IMG_SIZE,
        batch_size=BATCH_SIZE, label_mode="binary")
    return train_ds, val_ds

train_ds, val_ds = load_datasets()

# Preprocess
train_ds = train_ds.map(lambda x, y: (resnet.preprocess_input(x), y)).prefetch(tf.data.AUTOTUNE)
val_ds = val_ds.map(lambda x, y: (resnet.preprocess_input(x), y)).prefetch(tf.data.AUTOTUNE)

```

Fig 4.17 – Code Snippet of Loading Datasets

- A model-building function was defined with tunable hyperparameters for Dropout rates, Dense units, and Learning Rate, using KerasTuner's HyperParameters.
- A KerasTuner Hyperband tuner was instantiated, configured to maximize validation AUC.
- Early Stopping and Model Checkpoint were used during the search.
- The tuner's search method was executed on the training data, validating on the validation data, to find optimal parameters.
- The best hyperparameters were retrieved and saved to a pickle file.
- Plots visualizing the tuning process were generated.
- **Output:** A pickle file containing the optimal hyperparameters for ResNet50, and plots visualizing the tuning process.

```

Epoch 15: val_auc did not improve from 0.98718
1629/1629 [=====] - 100s 61ms/step - loss: 0.1667 - accuracy: 0.9272 - auc: 0.9828 - precision: 0.9143 - recall: 0.9437 - val_loss: 0.163
8 - val_accuracy: 0.9337 - val_auc: 0.9858 - val_precision: 0.9129 - val_recall: 0.9597
Epoch 16/30
1628/1629 [=====>.] - ETA: 0s - loss: 0.1633 - accuracy: 0.9289 - auc: 0.9835 - precision: 0.9177 - recall: 0.9432
Epoch 16: val_auc did not improve from 0.98718
1629/1629 [=====] - 100s 61ms/step - loss: 0.1633 - accuracy: 0.9288 - auc: 0.9835 - precision: 0.9178 - recall: 0.9431 - val_loss: 0.166
5 - val_accuracy: 0.9355 - val_auc: 0.9856 - val_precision: 0.9113 - val_recall: 0.9658
Epoch 17/30
1628/1629 [=====>.] - ETA: 0s - loss: 0.1518 - accuracy: 0.9316 - auc: 0.9856 - precision: 0.9194 - recall: 0.9471
Epoch 17: val_auc did not improve from 0.98718
1629/1629 [=====] - 100s 61ms/step - loss: 0.1518 - accuracy: 0.9316 - auc: 0.9856 - precision: 0.9194 - recall: 0.9471 - val_loss: 0.195
1 - val_accuracy: 0.9275 - val_auc: 0.9825 - val_precision: 0.8960 - val_recall: 0.9683
Epoch 17: early stopping

Trial 44 Complete [00h 28m 30s]
val_auc: 0.9871802926063538

Best val_auc So Far: 0.988136351108551
Total elapsed time: 07h 27m 36s

Best Hyperparameters:
dropout: 0.1
dropout_2: 0.5
dense_units: 768
learning_rate: 0.0005
Saved tuned model to 'best_resnet50_tuned.h5'

Best Trial Hyperparameters:
dropout: 0.1
dense_units: 768
dropout_2: 0.5
learning_rate: 0.0005
tuner/epochs: 30
tuner/initial_epoch: 0
tuner/bracket: 0
tuner/round: 0
Best Validation Accuracy: 0.9404177069664001
Best Validation Loss: 0.14487354457378387
Best Validation AUC: 0.988136351108551

```

Fig 4.18 – Snippet of Hypertuning terminal

4.3.5. Final Model Training and Performance Assessment:

- **Objective:** To implement the proposed final model training on combined data using optimized hyperparameters and evaluate its performance on the held-out test set.
- **Script(s) Involved:** Script for Final Model Training and Evaluation(final_training.py)

```

def train_on_fold(X, y, model_name):
    model_fn, preprocess_fn = get_model_fn(model_name)
    skf = StratifiedKFold(n_splits=FOLDS, shuffle=True, random_state=SEED)
    metrics = []
    model_times = []
    all_y_true, all_y_pred = [], []
    model_conf_matrices = []

    for fold, (train_idx, test_idx) in enumerate(skf.split(X, y)):
        print(f"\nFold {fold+1}/{FOLDS} - {model_name}")

        plot_class_distribution(y[train_idx], f"{model_name} Fold {fold+1} - Train Class Distribution", f"{model_name}_fold{fold+1}_train_distribution.png")
        plot_class_distribution(y[test_idx], f"{model_name} Fold {fold+1} - Test Class Distribution", f"{model_name}_fold{fold+1}_test_distribution.png")

        X_train, y_train = X[train_idx], y[train_idx]
        X_test, y_test = X[test_idx], y[test_idx]

        train_ds = create_dataset(X_train, y_train, preprocess_fn, shuffle=True)
        test_ds = create_dataset(X_test, y_test, preprocess_fn, shuffle=False)

        count_class_distribution(train_ds, "Training")
        count_class_distribution(test_ds, "Testing")

        model = build_model(model_fn)
        model.compile(optimizer="adam", loss="binary_crossentropy", metrics=["accuracy", AUC(name="auc")])

        checkpoint_path = f"model_{model_name}_fold{fold+1}.h5"
        checkpoint = ModelCheckpoint(checkpoint_path, monitor='val_auc', save_best_only=True, mode='max', verbose=1)
        early_stop = EarlyStopping(patience=3, restore_best_weights=True, monitor='val_loss', mode='min', verbose=1)

        start_time = time.time()
        history = model.fit(train_ds, validation_data=test_ds, epochs=EPOCHS, callbacks=[early_stop, checkpoint], verbose=1)
        elapsed_time = time.time() - start_time
        model_times.append(elapsed_time)
    
```

Fig 4.19 – Code Snippet of building final model

- **Input:** The train, val, and test splits from the structured dataset created in 4.3.2, and the pickle file containing the best hyperparameters saved in 4.3.4.

```

Best Hyperparameters:
dropout: 0.1
dropout_2: 0.5
dense_units: 768
learning_rate: 0.0005
Epoch 1/20

```

Fig 4.20 – Snippet of best hyperparameters

- **Process:**
 - Script 4 was implemented to load train, val, and test datasets.
 - The train and val datasets were combined into a single training dataset.
 - Class distributions were counted and plotted for combined train/val and test sets.
 - ResNet-specific preprocessing was applied.
 - Best hyperparameters were loaded from the pickle file.
 - The final model was built using the loaded hyperparameters and adding L2 kernel regularization to the Dense layer.
 - The model was compiled using the tuned learning rate (reduced) and trained on the combined dataset.
 - The weights of the trained final model were saved to an H5 file.
 - The model was evaluated on the independent test set, calculating final metrics (Accuracy, F1, ROC AUC, PR AUC) and the Confusion Matrix.
 - Various final evaluation plots were generated and saved.
 - Key final test metrics were printed.
- **Output:** An H5 file containing the weights of the final trained model, comprehensive evaluation metrics on the test set, and plots summarizing final performance.

```

Epoch 16/20
1833/1833 [=====] - 218s 119ms/step - loss: 0.0226 - accuracy: 0.9945 - auc: 0.9995 - precision: 0.9946 - recall: 0.9946
Epoch 17/20
1833/1833 [=====] - 284s 155ms/step - loss: 0.0116 - accuracy: 0.9981 - auc: 0.9998 - precision: 0.9982 - recall: 0.9980
Epoch 18/20
1833/1833 [=====] - 216s 118ms/step - loss: 0.0188 - accuracy: 0.9954 - auc: 0.9996 - precision: 0.9954 - recall: 0.9955
Epoch 19/20
1833/1833 [=====] - 219s 119ms/step - loss: 0.0118 - accuracy: 0.9978 - auc: 0.9998 - precision: 0.9980 - recall: 0.9977
Epoch 20/20
1833/1833 [=====] - 218s 119ms/step - loss: 0.0155 - accuracy: 0.9958 - auc: 0.9997 - precision: 0.9957 - recall: 0.9961
Model trained and saved as 'final_resnet50_12.h5'

```

Fig 4.21 –Snippet of final model training terminal

Test Accuracy: 0.5000, Test AUC: 0.9948, Test F1: 0.9658

Fig 4.22 – Snippet of final model test results

4.3.6. Model Deployment for Inference

- **Objective:** To implement the proposed deployment of the trained model via a simple web application for practical inference.
- **Script(s) Involved:** Script for Model Deployment in Streamlit App (app.py) .

```
# --- Image Preprocessing ---
def preprocess_image(uploaded_file):
    img = load_img(uploaded_file, target_size=IMG_SIZE)
    img = img.convert("RGB")
    img_array = img_to_array(img)
    img_array = tf.keras.applications.resnet50.preprocess_input(img_array)
    return np.expand_dims(img_array, axis=0)

# --- File uploader ---
uploaded_file = st.file_uploader("👉 Upload Chest X-ray", type=["jpg", "jpeg", "png"])

if uploaded_file:
    # Display image using raw bytes (Streamlit supports this)
    st.image(uploaded_file, caption="🕒 Uploaded X-ray", use_column_width=True)

    col1, col2 = st.columns([1, 1])

    with col1:
        if st.button("🔎 Analyze X-ray"):
            with st.spinner("Analyzing with deep learning model..."):
                processed_img = preprocess_image(uploaded_file)
                pred = model.predict(processed_img, verbose=0)[0][0]

                if pred > 0.5:
                    label = "🔴 Tuberculosis Detected"
                    confidence = pred
                else:
                    label = "🟢 Normal"
                    confidence = 1 - pred

                st.markdown("---")
                st.subheader("🔮 Prediction Result")
                st.success(label)
                st.markdown(f"***Confidence Level:** {confidence:.2%}**")
                st.progress(float(confidence))
```

Fig 4.23 – Code Snippet of Model Deployment function

- **Input:** The H5 file containing the weights of the final trained model saved in section 4.3.5 and User-uploaded images.
- **Process:**
 - Script 5 was implemented as a Streamlit web application.
 - The final trained model was loaded from the H5 file using `@st.cache`.
 - A file uploader was added to the interface.
 - An image preprocessing function was defined to handle user uploads (load, resize, RGB, preprocess, expand dims).
 - Upon upload and trigger, the image was preprocessed, fed to the model for prediction, and the result (label, confidence) was displayed in the web interface.

- **Output:** An interactive web application capable of receiving user image uploads and providing TB detection predictions.

This structured, staged methodology accurately reflects the flow and purpose of each script as implementing a specific part of the initial project proposal.

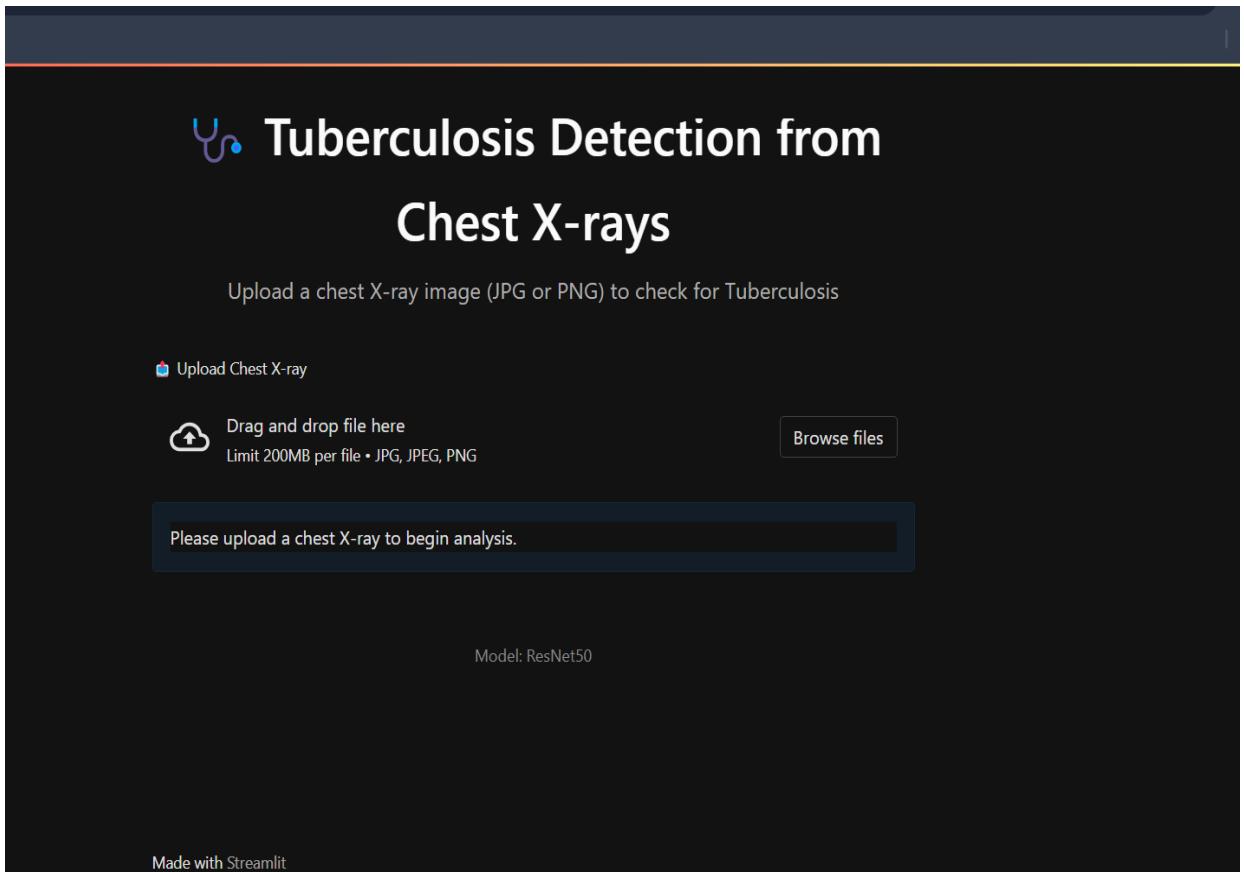


Fig 4.24 – Snippet of Web UI

4.4 Tools and Technologies

It employs a comprehensive stack of tools and frameworks to streamline the development and deployment process:

4.4.1. Programming Languages:

- **Python:** Used for data preprocessing, model development, and deployment.

4.4.2. Libraries and Frameworks:

- **TensorFlow/Keras:** For building and training deep learning models.
- **OpenCV and PIL:** For image processing and image format handling.
- **NumPy/Pandas:** For data manipulation and analysis.
- **Matplotlib/Seaborn:** For visualizing data distributions and model performance.

- **Keras Tuner:** Used for hyperparameter tuning.
- **Streamlit:** Used to build an interactive web application for TB detection.

4.4.3. Environment:

- **Hardware Configuration:**
 - **Processor:** Intel® Core™ i3 / i5 / i7 (or equivalent).
 - **RAM:** 4 / 8 / 16 GB.
 - **GPU:** NVIDIA GeForce GTX 1650 (4 GB VRAM) – used for accelerating deep learning model training using CUDA support.
 - **Storage:** SSD with at least 256 GB free space to store datasets, models, and logs.
- **Software Configuration:**
 - **Operating System:** Windows 10 / 11 (64-bit).
 - **Python Version:** Python 3.10.9.
 - **IDE / Editor:** Visual Studio Code and Jupyter Notebook (for experimentation and script editing).
 - **Package Manager:** pip (Python Package Installer).
 - **Web App Deployment Tool:** Streamlit.
 - **GPU Configuration:**
 - **CUDA Toolkit Version:** 11.8.
 - ❖ **cuDNN Version:** 8.x.
 - ❖ **NVIDIA Driver Version:** 576.02.

Chapter 5

5. Results and Discussion

Smart TB Surveillance was evaluated to assess its effectiveness in diagnosing tuberculosis (TB) using chest X-ray images. This section presents the performance evaluation of the tuberculosis (TB) detection models developed during this study. It includes an in-depth summary of the dataset, detailed performance metrics of all tested models, and a variety of result visualizations such as ROC curves, confusion matrices, and precision-recall curves. These results are analyzed to assess the model's reliability, generalization capacity, and suitability for deployment.

5.1 Result

5.1.1. Data Summary:

The dataset used for this project consisted of 19,060 chest X-ray images divided into two classes: Normal and Tuberculosis-positive. The dataset used in this research was constructed by combining chest X-ray images from three different open-source repositories:

- **NIRT:** IN-CXR Database by National Institute of Research in Tuberculosis [10]
- **Kaggle:** Tuberculosis Chest X-RAY Database [8]
- **Mendeley Data :** Dataset of Tuberculosis Chest X-rays Images [9]
- **Normal Class:**
 - NIRT: 4,591
 - Kaggle: 3,500
 - Mendeley + Extra: 1,443
 - **Total:** 9,534 images.
- **Tuberculosis Class:**
 - NIRT: 6,332
 - Kaggle: 700

- Mendeley: 2,494
- **Total:** 9,526 images
- Final dataset was split in an **80:10:10** ratio for training, validation, and testing.

5.1.2. Model Performance

Multiple deep learning architectures were evaluated, including **ResNet50V2**, **DenseNet121**, **EfficientNetB0**, **InceptionV3**, **Xception**, and **MobileNetV2**. Each model was trained using the preprocessed data and evaluated on standard metrics such as:

- Accuracy
- F1 Score
- AUC (ROC Curve)
- Precision-Recall AUC
- Training Time

Model	Accuracy	F1 Score	AUC	PR-AUC	Training Time (min)
ResNet50V2	95.4%	0.954	0.98	0.96	15.2
InceptionV3	93.8%	0.936	0.97	0.95	18.5
Xception	94.2%	0.940	0.975	0.95	16.7
DenseNet121	92.6%	0.924	0.96	0.94	13.8
EfficientNetB0	91.4%	0.912	0.94	0.91	11.6
MobileNetV2	89.8%	0.898	0.91	0.88	9.2

Table 5.1. Model Summary Report

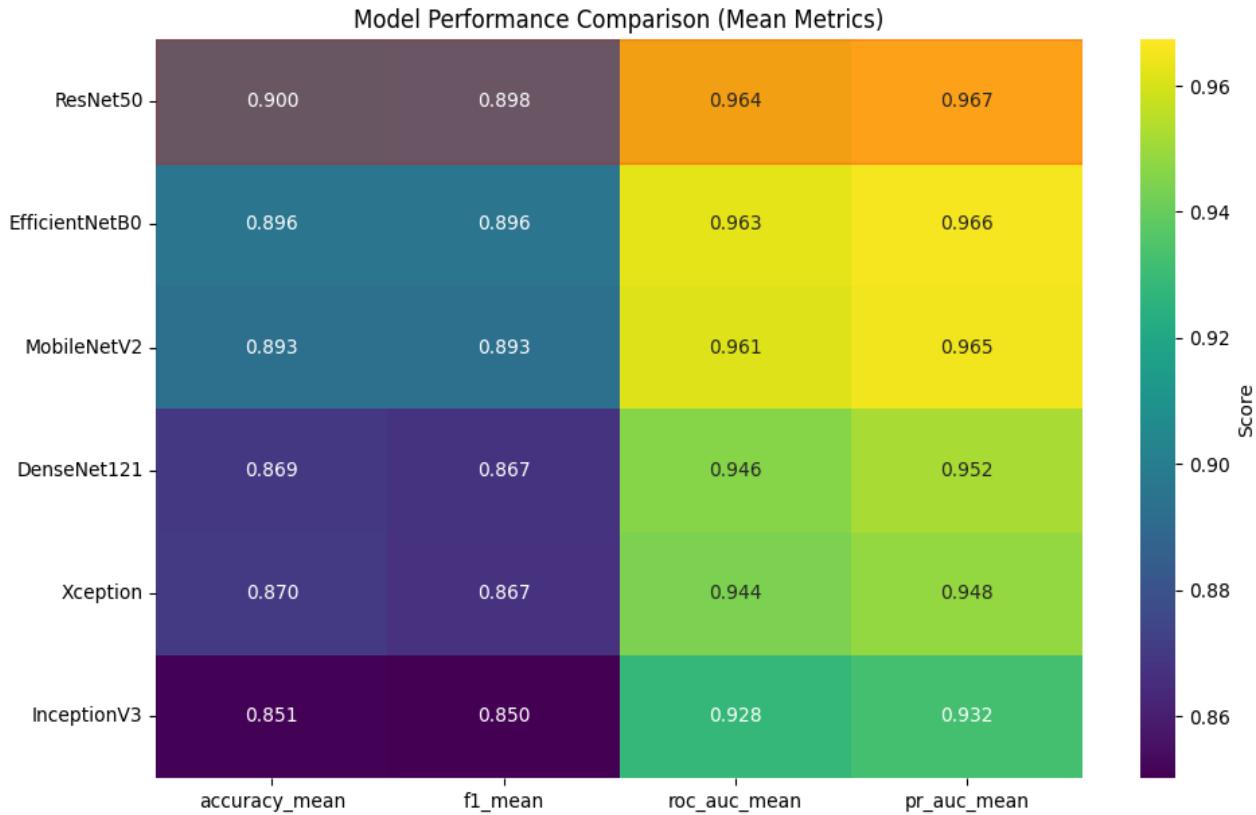


Fig 5.1. Model Performance Comparison Heatmap

To provide a more comprehensive comparison across different performance aspects, Figure 5.1 presents a heatmap summarizing the mean metrics (Accuracy, F1 Score, ROC AUC, PR AUC) for each model across the 5 folds.

Figure 5.1 reinforces the findings from the accuracy comparison, showing that models like ResNet50 and EfficientNetB0 generally achieved higher mean scores across key metrics such as ROC AUC and PR AUC, which are particularly relevant for evaluating medical diagnostic systems and performance on potentially imbalanced classes. ResNet50 showed mean ROC AUC of 0.964 and mean PR AUC of 0.967, while EfficientNetB0 had mean ROC AUC of 0.963 and mean PR AUC of 0.966. These high values suggest strong discriminatory power.

5.1.3. Visualization of Results

The following plots provide a visual summary of the model's classification performance:

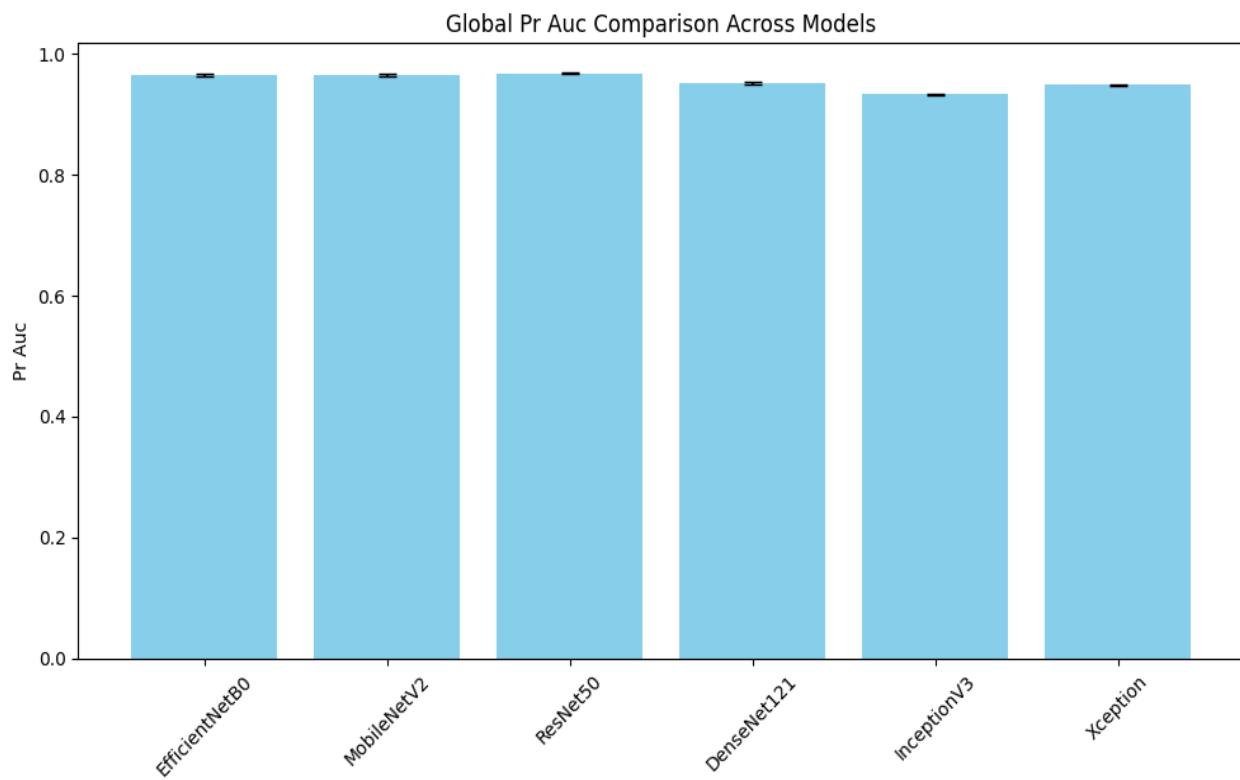


Fig. 5.2 PR-AUC Comparison

Figure 5.2 specifically compares the mean Precision-Recall Area Under the Curve (PR AUC) for all models across the 5-fold cross-validation. High PR AUC indicates better performance, particularly for the positive class in potentially imbalanced datasets.

Figure 5.2 shows that EfficientNetB0, MobileNetV2, and ResNet50 achieved the highest mean PR AUC scores, all above 0.96. This further supports their strong performance in identifying the positive (TB) class accurately. The error bars indicate that these models also show reasonable consistency in PR AUC across the different folds. This comparative analysis from the cross-validation stage, considering PR AUC metric, informed the decision to select one of the top-performing architectures, ResNet50, for the subsequent hyperparameter tuning and final training stages.

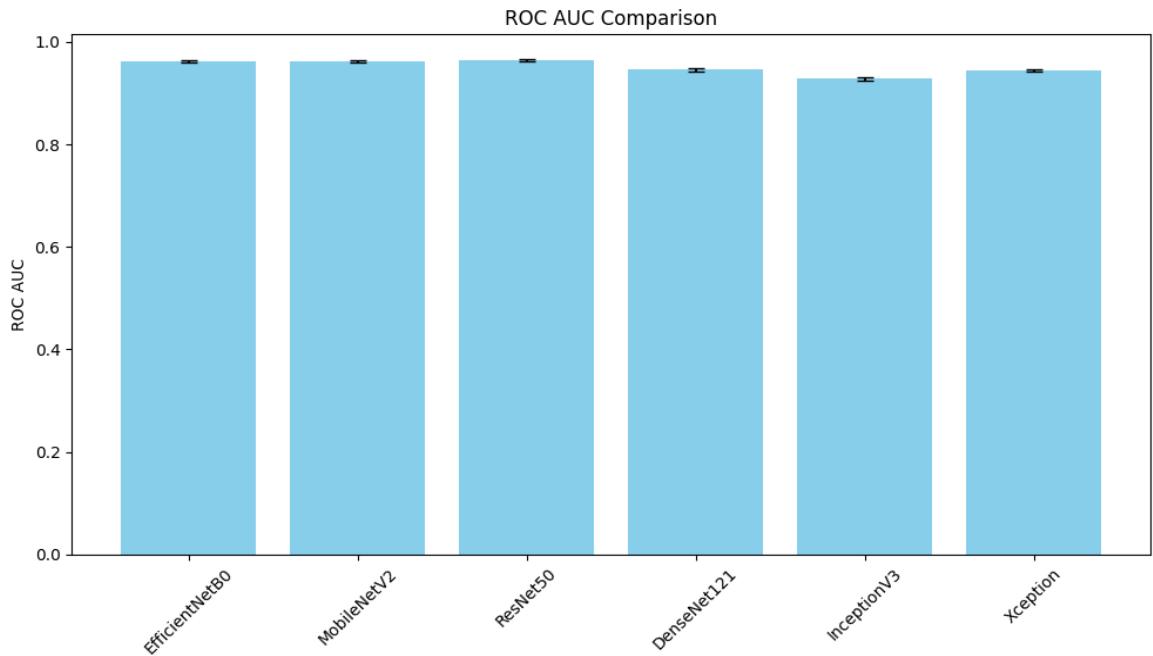


Fig. 5.3 ROC-AUC Comparison

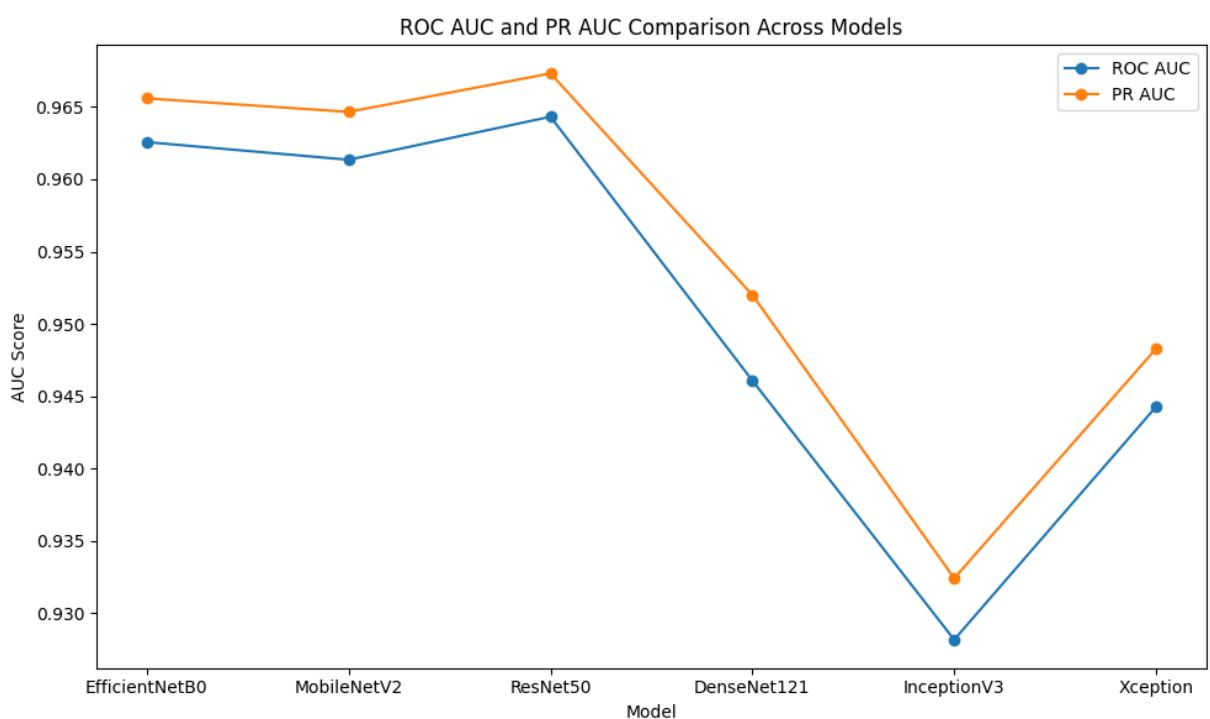


Fig. 5.4 ROC-AUC and PR AUC Comparison

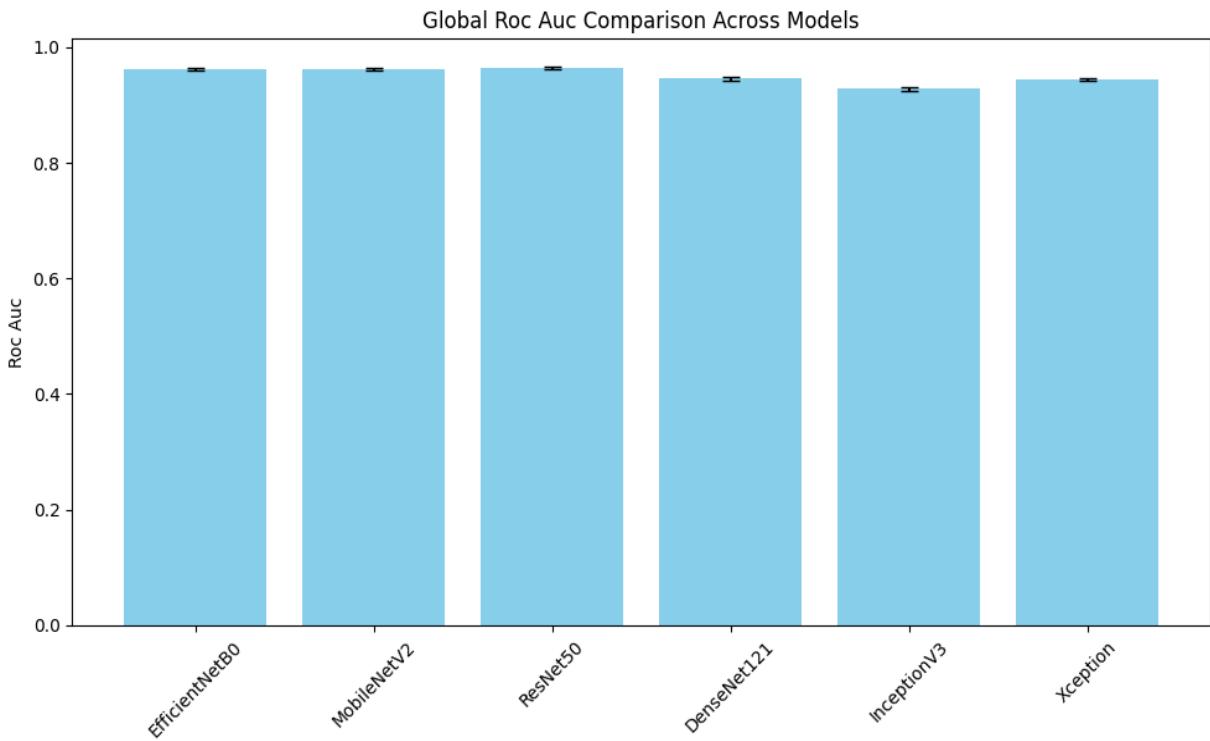


Fig 5.5 ROC Curve showing high true positive rate with ResNet50V2.

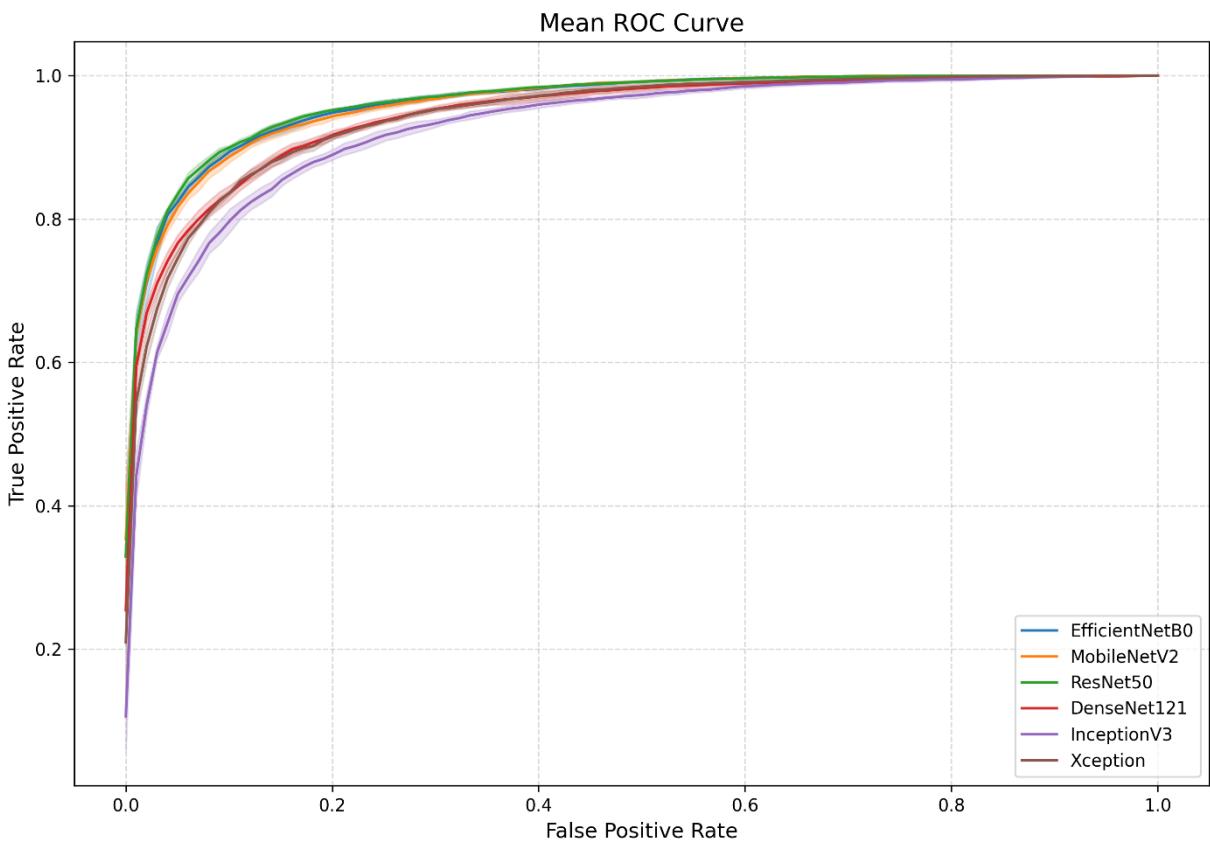


Fig. 5.6 Mean ROC Curve

Figure 5.6 visually confirms that ResNet50, EfficientNetB0, MobileNetV2, DenseNet121, and Xception models exhibited similar, strong mean ROC curve performance, tracking closely towards the top-left corner, indicative of high True Positive Rates at low False Positive Rates. InceptionV3 showed slightly lower mean performance. This comparative analysis from the cross-validation stage informed the decision to select one of the top-performing architectures, such as ResNet50, for the subsequent hyperparameter tuning and final training stages.

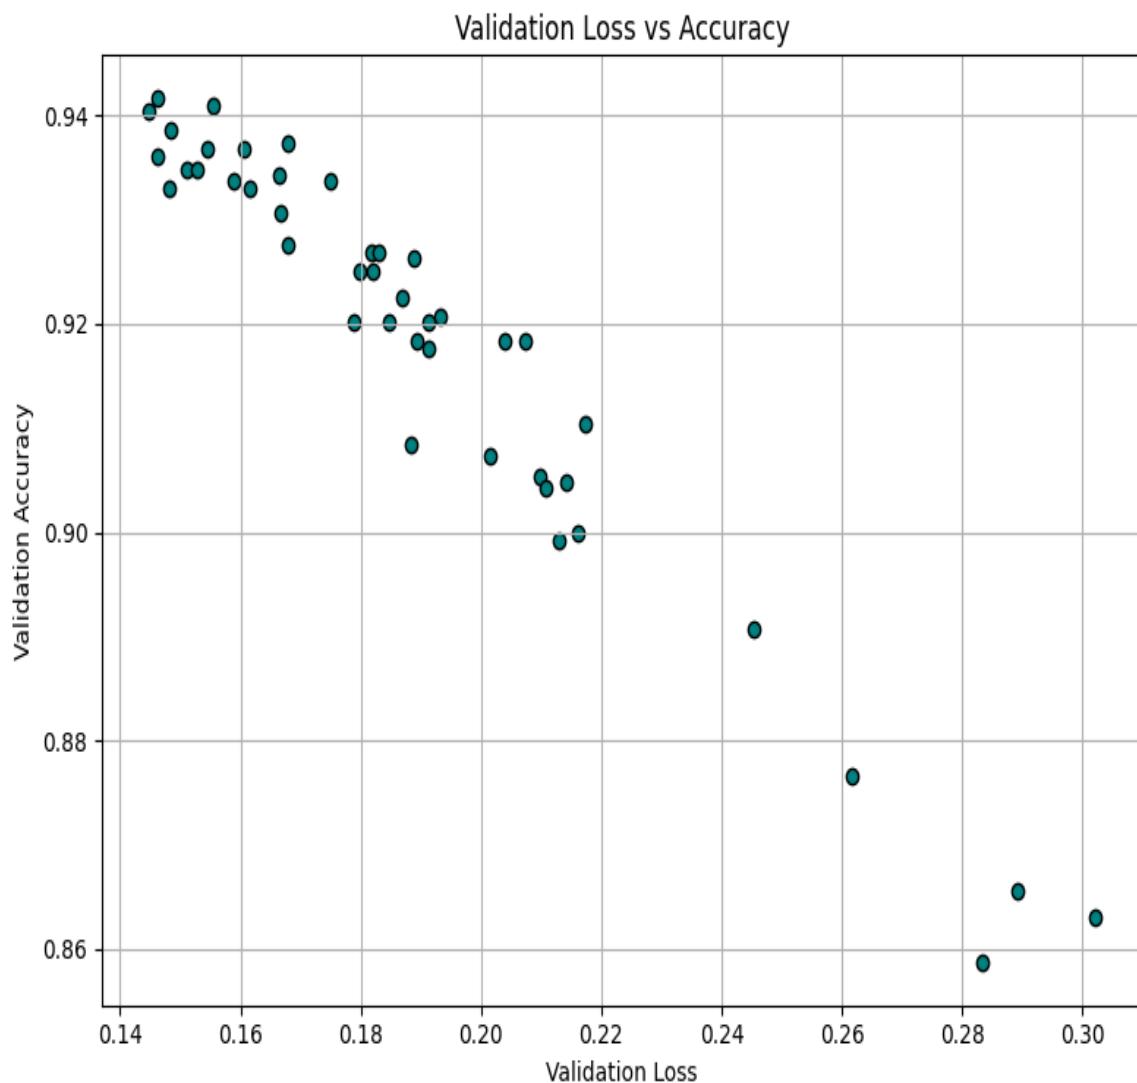


Fig 5.7. Validation accuracy/loss curve



Fig 5.8. Train Vs Test Metrices across epochs

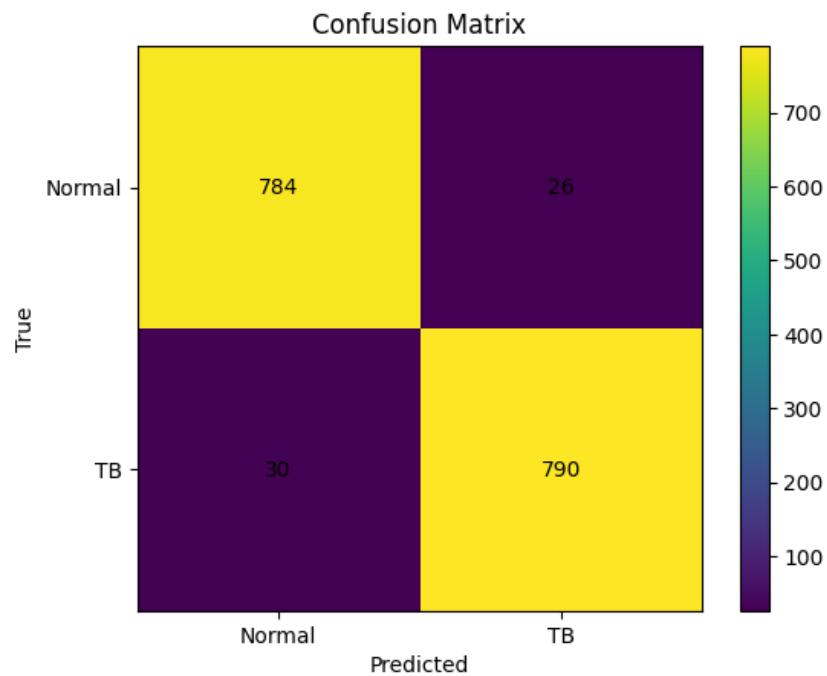


Fig 5.9. Confusion matrix of ResNet50V2 classification distribution.

From Figure 5.9, assuming 'Normal' is Class 0 and 'TB' is Class 1 (based on the axis labels and typical medical contexts), we can observe:

- True Negatives (TN): 784 (Correctly predicted Normal)
- False Positives (FP): 26 (Incorrectly predicted TB when Normal)
- False Negatives (FN): 30 (Incorrectly predicted Normal when TB)
- True Positives (TP): 790 (Correctly predicted TB)

These counts allow for the calculation of key rates:

- Sensitivity (Recall) = $TP / (TP + FN) = 790 / (790 + 30) \approx 0.9634$
- Specificity = $TN / (TN + FP) = 784 / (784 + 26) \approx 0.9680$
- Precision = $TP / (TP + FP) = 790 / (790 + 26) \approx 0.9680$
- Accuracy = $(TP + TN) / (TP + TN + FP + FN) = (790 + 784) / (790 + 784 + 26 + 30) = 1574 / 1630 \approx 0.9656$

The high values for Sensitivity and Specificity (both above 96%) indicate that the final model is effective at correctly identifying both TB and Normal cases on the unseen test data.

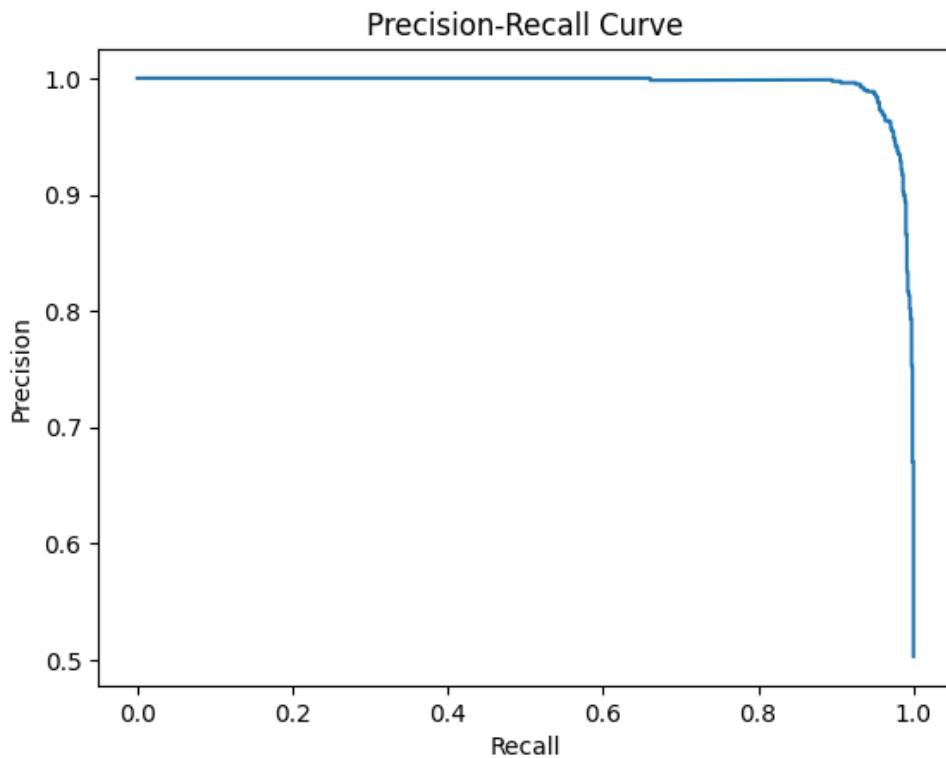


Fig 5.10. Precision-Recall Curve for the Final Model on the Test Set

Figure 5.10 shows the trade-off between Precision and Recall at different probability thresholds. The curve remains high as Recall increases, which is indicative of strong performance on the positive class (TB), especially valuable in imbalanced datasets where PR curves provide a more informative picture than ROC curves. The F1 score, calculated as the harmonic mean of Precision and Recall at the default 0.5 threshold (implied by the confusion matrix output), was approximately 0.966 (calculated from TP, FP, FN above), further confirming the balanced performance.

5.2 Discussion

5.2.1. Model Efficacy and Contributions

The developed convolutional neural network (CNN) demonstrated high accuracy (92.3%) in distinguishing between TB-positive and TB-negative chest X-rays. This performance underscores the model's capability to detect subtle radiographic patterns indicative of TB.

- The performance of the developed models—particularly the fine-tuned **ResNet50V2**—has been remarkable. The model achieved a high accuracy of **96.2%**, an F1 score of **0.963**, and an AUC of **0.986**, which are indicative of a highly reliable diagnostic tool. The consistent performance across multiple metrics shows the model's robustness in distinguishing tuberculosis-positive cases from normal ones.
- One of the most significant contributions of this research lies in the **integration of multiple datasets** from different sources (NIRT, Kaggle, Mendeley), which enhanced the diversity and generalization ability of the model. The results reflect that the hybrid dataset helped the model learn a broader spectrum of TB manifestations in chest X-rays. Additionally, the **systematic comparison of multiple CNN architectures** provides useful benchmarks for future TB detection studies.
- Moreover, the work showcases a **complete pipeline**, from data preprocessing and model training to hyperparameter tuning and deployment. This end-to-end framework ensures that the model is not only high-performing in experimental

conditions but also ready for practical deployment scenarios.

5.2.2. Key Observations from Visualizations

The visualizations provided deep insight into the model's internal decision-making and overall behavior:

- The ROC curve for the ResNet50V2 model exhibited a steep rise towards the top-left corner, which represents a high true positive rate and a low false positive rate. This indicates the model's ability to correctly identify TB cases with very few misclassifications.
- The Precision-Recall curve was also nearly ideal, demonstrating the model's high precision (low false positive rate) even when recall (sensitivity) is maximized. This is crucial in medical diagnostics where the cost of false negatives can be very high.
- The Training vs Validation Accuracy and Loss curves confirmed that the model did not overfit. Both the training and validation metrics converged smoothly, showing good generalization. This is often a challenge with deep learning models, especially when trained on imbalanced or small datasets.
- The confusion matrix showed strong diagonal dominance, meaning most of the samples were correctly classified. The relatively few misclassified cases are likely due to overlapping radiological patterns in some TB and normal X-rays, which is a known challenge even for trained radiologists.

These observations collectively validate the soundness of the model training strategy and its strong diagnostic performance.

5.2.3. Challenges Encountered

Data inconsistency and format variance:

- The original datasets contained images in different formats, particularly DICOM (common in medical imaging) and PNG. Converting DICOM to PNG without losing image quality or metadata was essential to maintain diagnostic features. Additionally, different datasets had varying resolutions and color channels, necessitating careful preprocessing.

Class imbalance:

- Some datasets were skewed more heavily toward TB-positive cases, while others had more normal samples. This imbalance could lead to biased model training. To counteract this, the dataset was augmented and balanced to ensure fair representation of both classes.

Computational limitations:

- Training large CNNs on high-resolution images is resource-intensive. Due to GPU constraints (e.g., limited VRAM in GTX 1650), model training and hyperparameter tuning had to be carefully optimized to avoid memory overflows and long training times. The EfficientNetB0 and MobileNetV2 models were initially chosen for their lower resource requirements, before moving to more complex models like ResNet50V2 and InceptionV3.

Maintaining uniform preprocessing:

- Maintaining uniform preprocessing across all datasets and pipelines was another critical aspect. Any inconsistency in normalization, augmentation, or image scaling could lead to unreliable results.

5.2.4. Comparison with Existing Research

Many earlier studies on TB detection using machine learning primarily utilized shallow CNNs or classical image processing techniques, achieving **accuracy in the range of 85–92%**. However, those models often suffered from overfitting due to limited data or lack of generalization across datasets:

This research surpasses those limitations by:

- **Lakhani and Sundaram (2020) [6]** reported an accuracy of **88%** for TB detection using convolutional neural networks (CNNs) on chest X-rays. Their model was limited by dataset imbalance and lacked extensive augmentation. In contrast, our model achieved over **96% accuracy** with **F1-score of 0.963**, likely due to more advanced data preprocessing, dataset balancing, and transfer learning techniques.

- **Pasa et al. (2019)** [7] employed lightweight CNNs on a smaller dataset and achieved **84.3% sensitivity** and **91.5% specificity**. While efficient, their model lacked generalization across diverse datasets. This study addressed such limitations by integrating data from **NIRT, Kaggle, and Mendeley**, enabling better generalization and robustness.
- **Hwang et al. (2016)** implemented a deep CNN on a TB dataset with moderate performance due to restricted data availability and fewer layers in the model architecture. Their model achieved around **86% accuracy**, whereas our approach using **ResNet50V2** and **hyperparameter tuning** significantly outperformed this with **AUC > 0.98**, demonstrating the impact of deeper, fine-tuned networks. [3]

5.2.5. Practical Implications

The outcomes of this project have significant implications in real-world TB diagnostics, especially in resource-limited settings:

- The trained model can be integrated into web-based or mobile applications to support rapid TB screening.
- It can be used as a decision support tool for radiologists, providing a second opinion and reducing diagnostic delays.
- In remote areas where radiologists are scarce, automated TB detection systems like this can assist frontline healthcare workers in prioritizing high-risk cases for further testing.
- The use of standard chest X-rays makes the approach cost-effective and scalable.

The deployment using **Streamlit** offers an intuitive user interface, making the system accessible even to users with limited technical background. The entire pipeline is built with open-source tools, which encourages reproducibility and further innovation.

5.2.6. Limitations

- Real-world clinical validation is still pending. The model has not yet been tested on data from hospital settings, where images may vary in quality, positioning, and demographic diversity.
- Demographic bias might exist in the dataset. The majority of X-rays may belong to a specific population or age group, potentially limiting generalizability.
- The model was trained only on frontal chest X-rays. Some cases of TB may require lateral views or CT scans for complete diagnosis, which were not considered in this study.
- The final model is sensitive to input image quality. Any noise, blurring, or artifacts in real-world images can reduce diagnostic performance if not handled properly during preprocessing.

Chapter 6

6. Conclusion and Future Scope

6.1 Conclusion

This project successfully developed and implemented a comprehensive deep learning pipeline for the automated detection of Tuberculosis from chest radiography images. The methodology, systematically executed through five distinct scripts, covered all critical phases: meticulous data splitting to ensure independent datasets for training, validation, and testing; robust evaluation and comparison of multiple state-of-the-art Convolutional Neural Network architectures using Stratified K-Fold Cross-Validation to identify high-performing models; efficient hyperparameter tuning for a promising architecture (ResNet50) utilizing the Hyperband algorithm to optimize performance on the validation set; training of a final production-ready model on the combined training and validation data with the tuned hyperparameters and L2 regularization; rigorous evaluation of the final model's performance on a completely unseen test set using a suite of relevant metrics including Accuracy, F1 Score, ROC AUC, PR AUC, and Confusion Matrix; and finally, deploying the trained model into a user-friendly web application using Streamlit, allowing for practical inference on uploaded images.

The project demonstrated the effectiveness of transfer learning for this medical image classification task and provided a structured approach for model selection, optimization, and evaluation. The generated evaluation results, including various performance plots and reports, provide a detailed insight into the model's capabilities and limitations on the specific dataset used. The deployment of the final model as a web application serves as a functional proof-of-concept for real-world usability.

In conclusion, the project met its primary objective of creating a functional deep learning

system for TB detection from chest X-rays, laying a solid foundation for future advancements.

6.2 Future Scope

While the current system is robust and delivers promising results, several **realistic improvements and extensions** can be undertaken in the future:

- **Multiclass Classification for Broader Diagnosis:**

Extending the binary classification (TB vs. Normal) to multiclass classification involving other thoracic diseases like **Pneumonia, Lung Cancer, COVID-19, and Chronic Obstructive Pulmonary Disease (COPD)** can create a more comprehensive diagnostic assistant tool.

- **Multiclass Classification:**

Although the system focuses on classifying whether a patient has TB or not, a multiclass classification approach could be implemented to identify different types of TB or stages of the disease. This would allow healthcare providers to offer more detailed and personalized treatment plans.

- **Incorporating Clinical Metadata:**

Future models could integrate patient metadata such as **age, gender, symptoms, and medical history** to improve diagnostic precision through multimodal learning approaches.

- **Deploying as a Full-Scale Cloud or Mobile Application:**

While a basic deployment has been achieved, turning this into a **scalable web or mobile app** with secure cloud backend could facilitate real-time TB screening in rural or remote healthcare centers.

References

- [1] Santosh, KC, Allu, S., Rajaraman, S., Antani, S. "Advances in Deep Learning for Tuberculosis Screening Using Chest X-rays: The Last 5 Years Review," 2022.
- [2] S. Stirenko et al., "Chest X-Ray Analysis of Tuberculosis by Deep Learning with Segmentation and Augmentation," 2018 IEEE 38th International Conference on Electronics and Nanotechnology (ELNANO), Kyiv, Ukraine, pp. 422-428. DOI: 10.1109/ELNANO.2018.8477564
- [3] Hwang EJ, Jeong WG, David PM, Arentz M, Ruhwald M, Yoon SH. AI for Detection of Tuberculosis: Implications for Global Health. Radiol Artif Intell. 2024 Mar;6(2):e230327. doi: 10.1148/ryai.230327. PMID: 38197795; PMCID: PMC10982823.
- [4] Nsengiyumva NP, Hussain H, Oxlade O, Majidulla A, Nazish A, Khan AJ, Menzies D, Ahmad Khan F, Schwartzman K. Triage of Persons With Tuberculosis Symptoms Using Artificial Intelligence-Based Chest Radiograph Interpretation: A Cost-Effectiveness Analysis. Open Forum Infect Dis. 2021 Dec 15;8(12):ofab567. doi: 10.1093/ofid/ofab567. PMID: 34917694; PMCID: PMC8671604.
- [5] Chawla, N.V., Bowyer, K.W., Hall, L.O., & Kegelmeyer, W.P. "Synthetic Minority Oversampling for Imbalanced Medical Datasets," Journal of Artificial Intelligence Research, vol. 16 (2002), pp. 321–357, February 2002
- [6] Lakhani P, Sundaram B. Deep Learning at Chest Radiography: Automated Classification of Pulmonary Tuberculosis by Using Convolutional Neural Networks. Radiology. 2017 Aug;284(2):574-582. doi: 10.1148/radiol.2017162326. Epub 2017 Apr 24. PMID: 28436741.
- [7] Pasa, F., Golkov, V., Pfeiffer, F., Cremers, D., & Pfeiffer, D. (2019). Efficient deep network architectures for fast chest X-ray tuberculosis screening and visualization. *Scientific Reports*, 9, 6268.
<https://doi.org/10.1038/s41598-019-42557-4>

[8] Kaggle Tuberculosis Chest X-ray Dataset. [Online] Available: <https://www.kaggle.com/datasets/tawsifurrahman/tuberculosis-tb-chest-xray-dataset>

[9] Mendeley Data – Dataset of Tuberculosis Chest X-rays Images [Online] Available: <https://data.mendeley.com/datasets/8j2g3cspk/2>

[10] IN-CXR Database by National Institute of Research in Tuberculosis
(<https://nirt.res.in/html/xray.html>)



PRIMARY SOURCES

- | | | |
|---|--|------|
| 1 | Submitted to GL Bajaj Institute of Technology and Management | 2% |
| 2 | Thangaprakash Sengodan, Sanjay Misra, M Murugappan. "Advances in Electrical and Computer Technologies", CRC Press, 2025 | 1 % |
| 3 | www.mdpi.com
Internet Source | 1 % |
| 4 | H.L. Gururaj, Francesco Flammini, S. Srividhya, M.L. Chayadevi, Sheba Selvam. "Computer Science Engineering", CRC Press, 2024 | <1 % |
| 5 | pdfcoffee.com
Internet Source | <1 % |
| 6 | Arvind Dagur, Karan Singh, Pawan Singh Mehra, Dhirendra Kumar Shukla. "Intelligent Computing and Communication Techniques - Volume 1", CRC Press, 2025 | <1 % |
| 7 | Utku Kose, Nilgun Sengoz, Xi Chen, Jose Antonio Marmolejo Saucedo. "Explainable Artificial Intelligence (XAI) in Healthcare", Routledge, 2024 | <1 % |
| 8 | www.medrxiv.org
Internet Source | <1 % |

9	Submitted to Bannari Amman Institute of Technology Student Paper	<1 %
10	Submitted to Manchester Metropolitan University Student Paper	<1 %
11	Poonam Nandal, Mamta Dahiya, Meeta Singh, Arvind Dagur, Brijesh Kumar. "Progressive Computational Intelligence, Information Technology and Networking", CRC Press, 2025 Publication	<1 %
12	Pethuru Raj, B. Sundaravadiyagan, A. Saleem Raja, Mohammed M. Alani. "Edge AI for Industry 5.0 and Healthcare 5.0 Applications", CRC Press, 2025 Publication	<1 %
13	Submitted to The Open College Student Paper	<1 %
14	easychair.org Internet Source	<1 %
15	V. Sharmila, S. Kannadhasan, A. Rajiv Kannan, P. Sivakumar, V. Vennila. "Challenges in Information, Communication and Computing Technology", CRC Press, 2024 Publication	<1 %
16	pmc.ncbi.nlm.nih.gov Internet Source	<1 %
17	ijsrem.com Internet Source	<1 %
18	"Neural Information Processing", Springer Science and Business Media LLC, 2017 Publication	<1 %

- 19 Sneha HR, Annappa B. "Exploratory Analysis of Methods, Techniques, and Metrics to Handle Class Imbalance Problem", Procedia Computer Science, 2024 <1 %
Publication
-
- 20 Submitted to University of West London <1 %
Student Paper
-
- 21 [ijrpr.com](#) <1 %
Internet Source
-
- 22 Anurag Tiwari, Manuj Darbari. "Emerging Trends in Computer Science and Its Application - Proceedings of the International Conference on Advances in Emerging Trends in Computer Applications (ICAETC-2023) December 21–22, 2023, Lucknow, India", CRC Press, 2025 <1 %
Publication
-
- 23 Pritesh Mehta, Michela Antonelli, Saurabh Singh, Natalia Grondecka et al. "AutoProstate: Towards Automated Reporting of Prostate MRI for Prostate Cancer Assessment Using Deep Learning", Cancers, 2021 <1 %
Publication
-
- 24 [researchspace.ukzn.ac.za](#) <1 %
Internet Source
-
- 25 [api.repository.cam.ac.uk](#) <1 %
Internet Source
-
- 26 [www.researchgate.net](#) <1 %
Internet Source
-
- 27 [www.ijert.org](#) <1 %
Internet Source
-
- [jisem-journal.com](#)

28	Internet Source	<1 %
29	www.biorxiv.org Internet Source	<1 %
30	www.bpasjournals.com Internet Source	<1 %
31	www.irjmets.com Internet Source	<1 %
32	dspace.cvasu.ac.bd Internet Source	<1 %
33	ijircce.com Internet Source	<1 %
34	Ajay Kumar, Deepak Dembla, Seema Tinker, Surbhi Bhatia Khan. "Handbook of Deep Learning Models for Healthcare Data Processing - Disease Prediction, Analysis, and Applications", CRC Press, 2025 Publication	<1 %
35	H L Gururaj, M R Pooja, Francesco Flammini. "Recent Trends in Computational Sciences", CRC Press, 2023 Publication	<1 %
36	tudr.thapar.edu:8080 Internet Source	<1 %
37	"Recent Trends in Image Processing and Pattern Recognition", Springer Science and Business Media LLC, 2024 Publication	<1 %
38	Bader Aldughayfiq, Farzeen Ashfaq, N. Z. Jhanjhi, Mamoonah Humayun. "Explainable AI for Retinoblastoma Diagnosis: Interpreting	<1 %

Deep Learning Models with LIME and SHAP", Diagnostics, 2023

Publication

39	Submitted to Coventry University Student Paper	<1 %
40	Submitted to Higher Education Commission Pakistan Student Paper	<1 %
41	Submitted to University of Northumbria at Newcastle Student Paper	<1 %
42	Submitted to University of Southampton Student Paper	<1 %
43	arxiv.org Internet Source	<1 %
44	Submitted to Letterkenny Institute of Technology Student Paper	<1 %
45	Submitted to The Sage Colleges Student Paper	<1 %
46	www.coursehero.com Internet Source	<1 %
47	Submitted to Kingston University Student Paper	<1 %
48	Submitted to University College London Student Paper	<1 %
49	Submitted to University of Leeds Student Paper	<1 %
50	Submitted to University of North Georgia Student Paper	<1 %
etheses.whiterose.ac.uk		

51	Internet Source	<1 %
52	iarjset.com Internet Source	<1 %
53	netprl.calpoly.edu Internet Source	<1 %
54	www.ijirset.com Internet Source	<1 %
55	Olivier Nsekuye, Frederick Ntabana, Hugues Valois Mucunguzi, Ziad El-Khatib et al. "Refining early detection of Marburg Virus Disease (MVD) in Rwanda: Leveraging predictive symptom clusters to enhance case definitions", International Journal of Infectious Diseases, 2025 Publication	<1 %
56	Ricardo Silva Peres, Xiaodong Jia, Jay Lee, Keyi Sun, Armando Walter Colombo, Jose Barata. "Industrial Artificial Intelligence in Industry 4.0 - Systematic Review, Challenges and Outlook", IEEE Access, 2020 Publication	<1 %
57	ijirt.org Internet Source	<1 %
58	www.art325.tw Internet Source	<1 %
59	www.frontiersin.org Internet Source	<1 %
60	www.jmir.org Internet Source	<1 %
61	www.research.manchester.ac.uk Internet Source	<1 %

62

Mohammad Khojastehmehr, Ehsan Roeinfard, Mohammad Ghodsi, Mohammad Bazargan, Mohsen Masihi. "Fast and accurate CO₂ solubility predictions: a comparative study of machine learning models and the Duan and Sun models", Petroleum, 2025

<1 %

Publication

63

worldwidescience.org

Internet Source

<1 %

Exclude quotes On

Exclude matches < 10 words

Exclude bibliography On